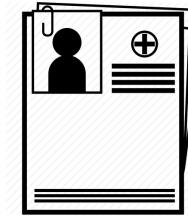
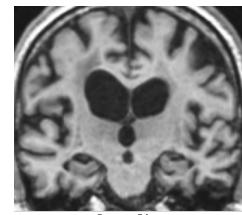
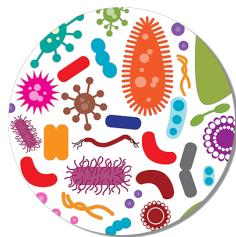
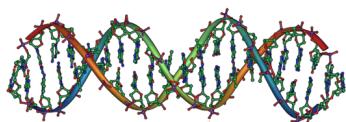
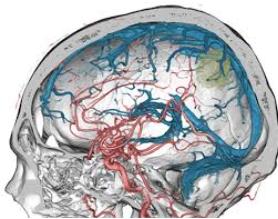


Latent variable models for the analysis of heterogeneous information

Marco Lorenzi

Epione Research Group, Université Côte d'Azur, Inria

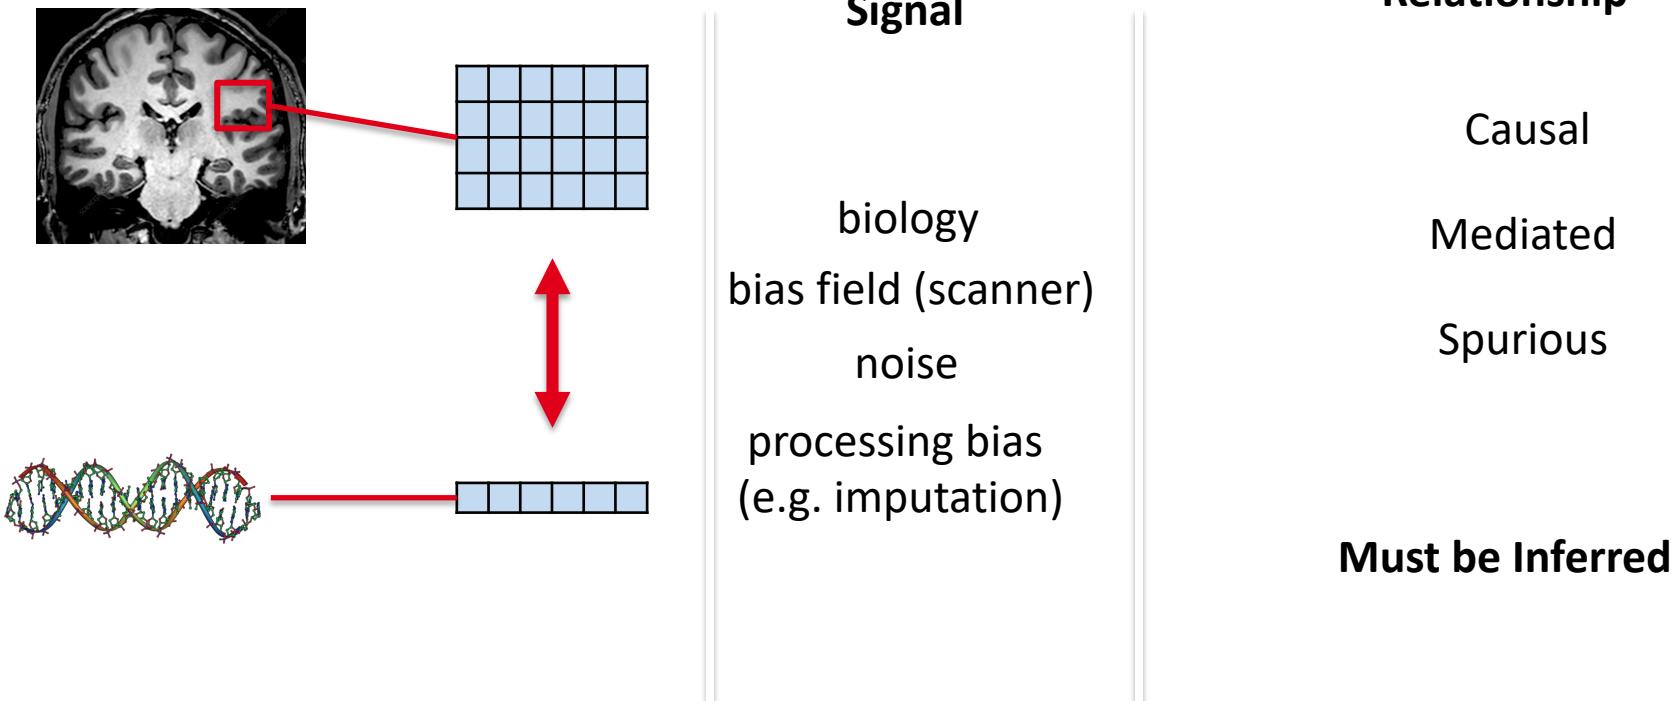
Data Science meets biomedical research



...

Variability of multivariate biomedical data

- within/between views -



Variability of multivariate biomedical data

- between datasets -



Population
Acquisition
Processing
Data security

...

Latent variable models

1. ***Multi-variate* modeling**
2. Novel scalable approaches to *multi-view* data

Latent variable models

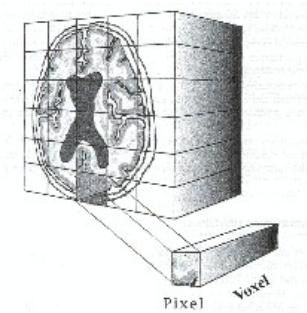
- 1. Multi-variate modeling**
2. Novel scalable approaches to multi-view data

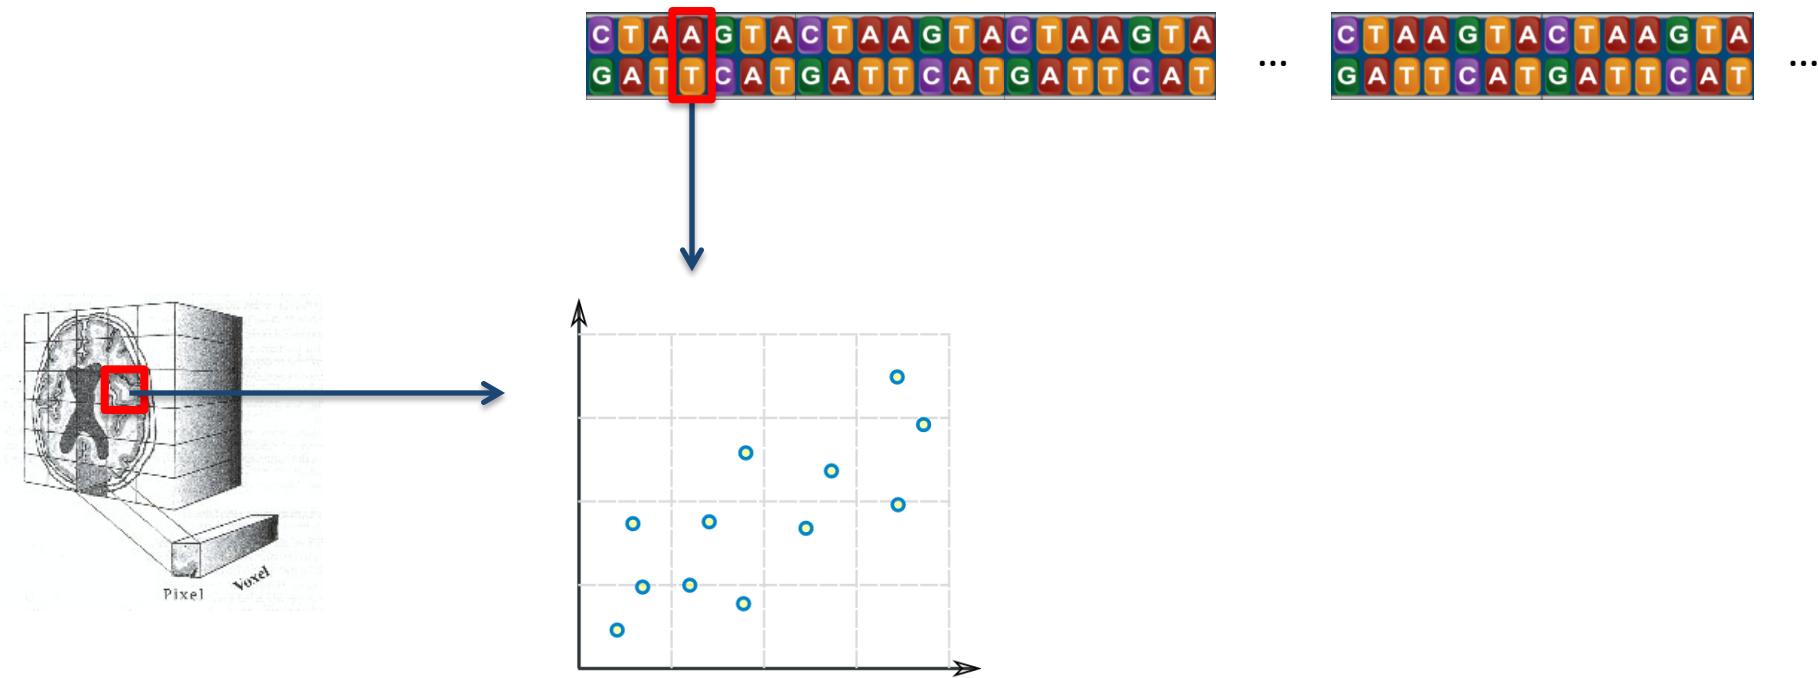
CTAAGTACTAAGTACTAAGTA
GATTTCATGATTTCATGATTTCAT

...

CTAAGTACTAAGTA
GATTTCATGATTTCAT

...



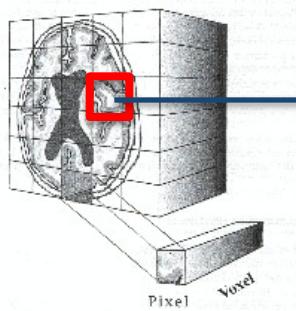
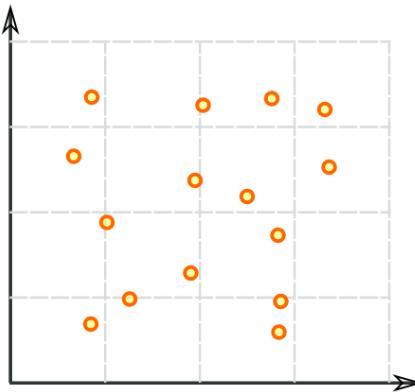


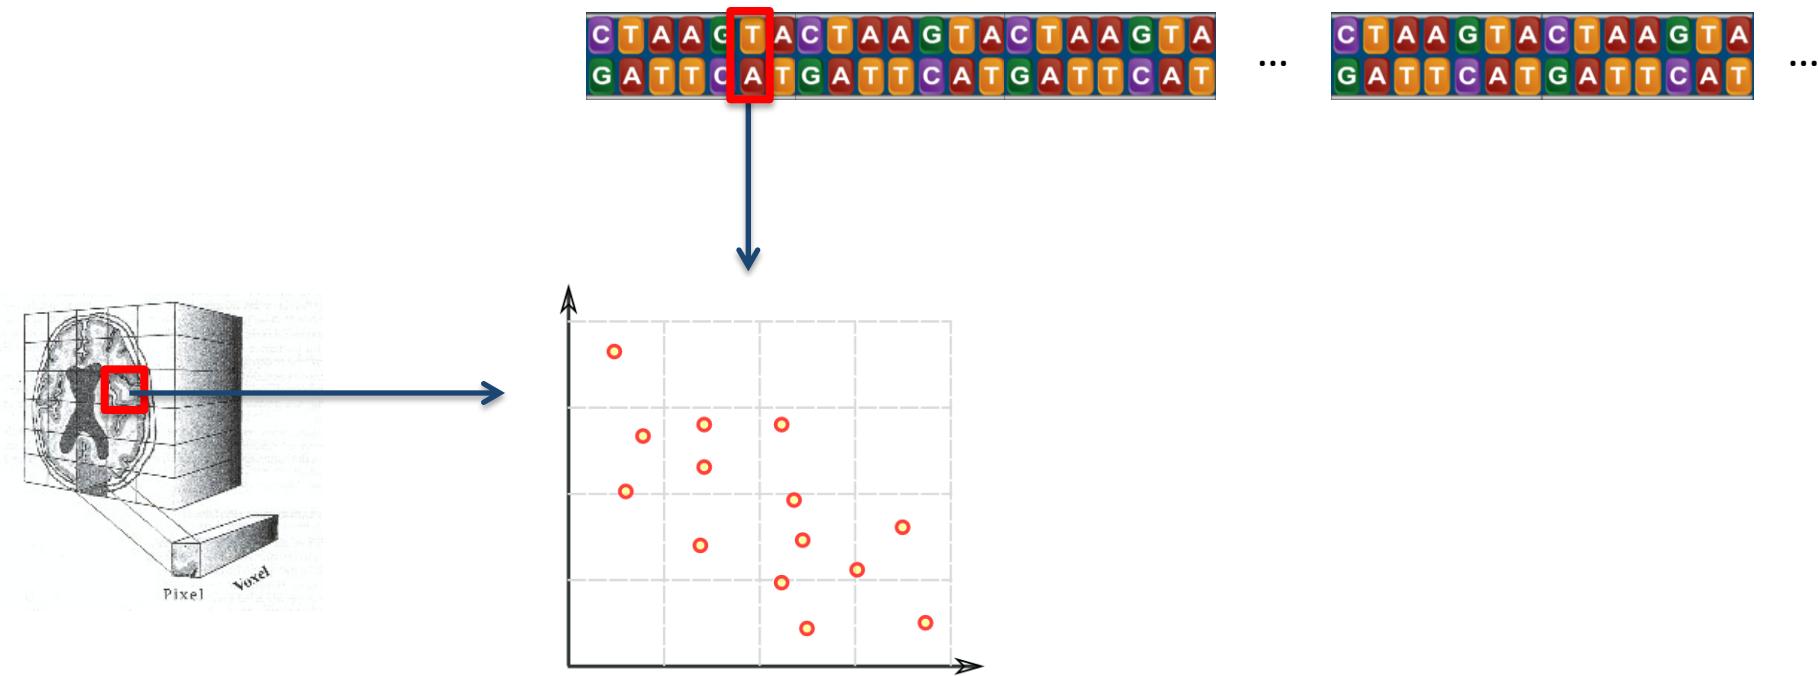
...

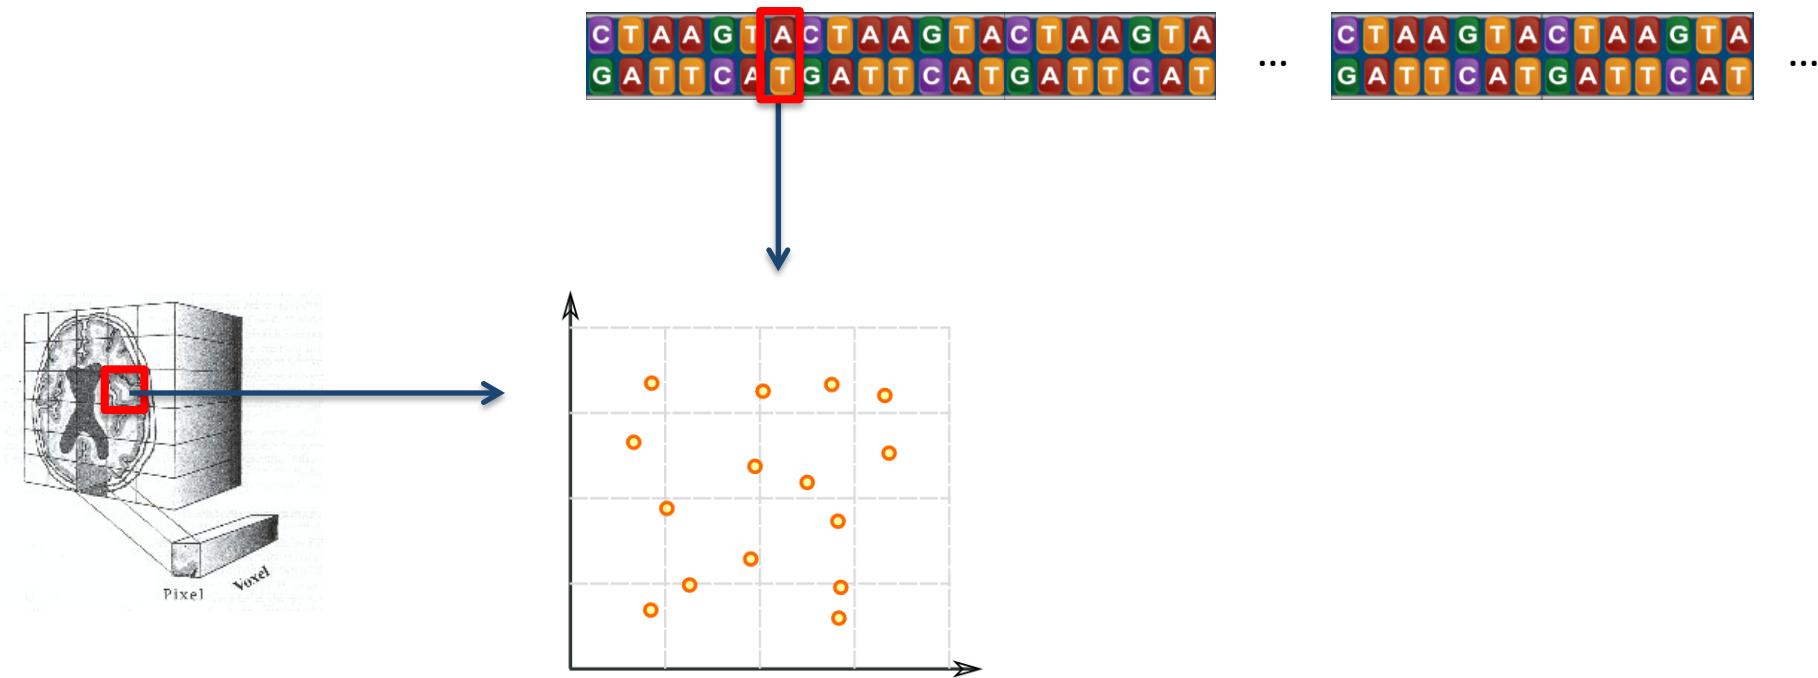
CTAAGTACTAAGTA
GATTCACTGATTCA

...

CTAAGTACTAAGTA
GATTCACTGATTCA

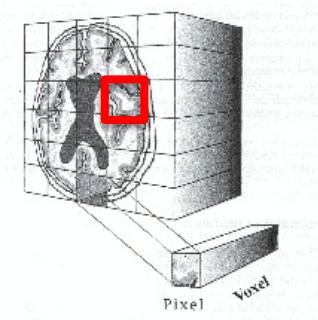






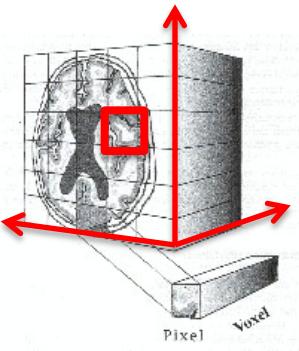


Iterate for > 1'000'000 variants





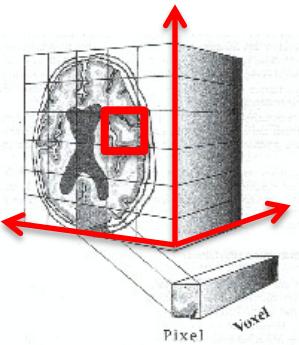
Iterate for > 1'000'000 variants



Iterate for > 1'000'000
image locations



Iterate for > 1'000'000 variants

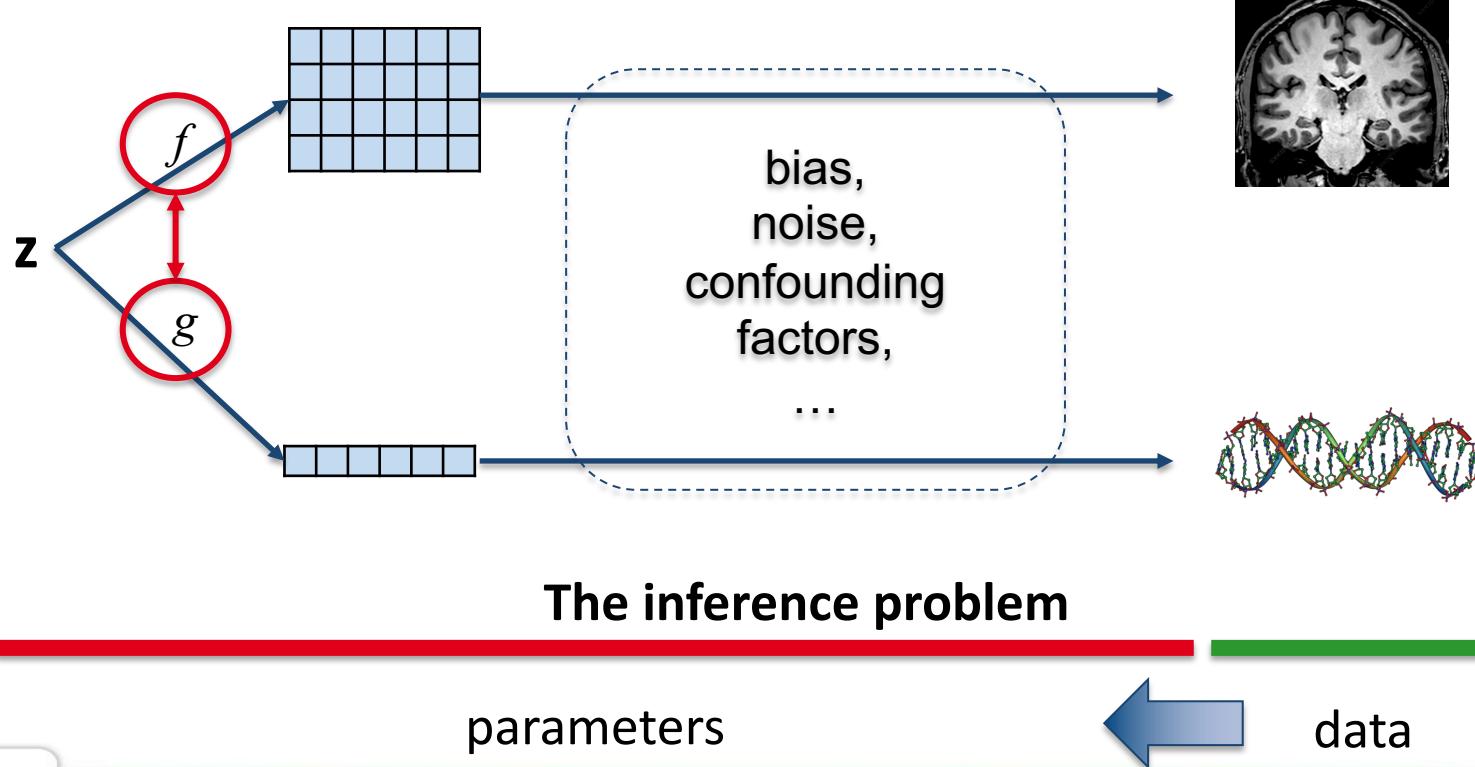


Iterate for > 1'000'000
image locations

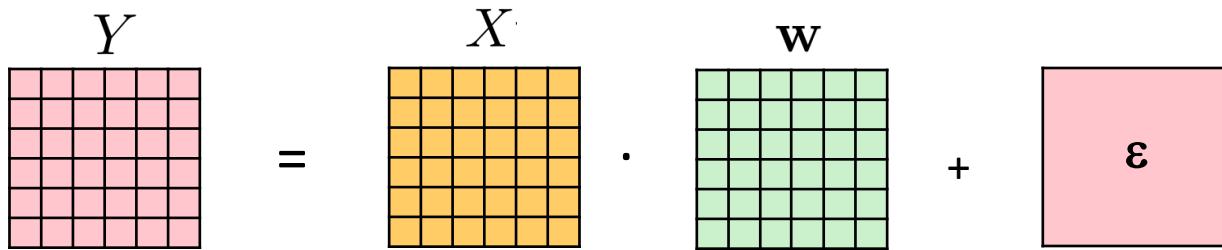
- Hard interpretability
- False positive discoveries
- No interaction across brain and genetic areas



A unified statistical formulation via latent generative models



The building-block: linear model

$$Y = X \cdot w + \varepsilon$$


$$\nabla_w (\|Y - Xw\|^2) = 0 \quad \longrightarrow \quad w = (X^T X)^{-1} X^T Y$$

Classical formulation of latent variable models

Principal component analysis

$$t \xrightarrow{\mathbf{w}^T} X$$
$$X = \begin{matrix} t \\ \vdots \end{matrix} \mathbf{w}^T + \varepsilon_X$$

inference

The diagram illustrates the classical formulation of latent variable models. On the left, a vector t is transformed by a weight matrix \mathbf{w}^T into a matrix X . On the right, X is shown as the sum of a deterministic component (t times \mathbf{w}^T) and a residual error ε_X . A dashed arrow labeled "inference" points from the residual error term towards the equation, indicating that the goal is to estimate the error term based on the observed data X .

Classical formulation of latent variable models

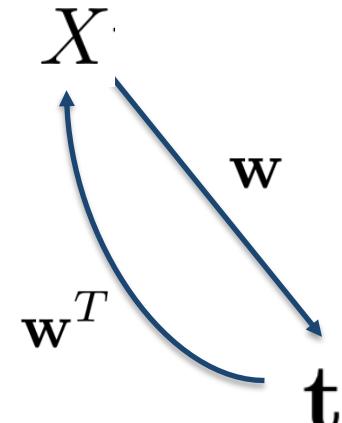
Principal component analysis

A **variance** maximisation problem:

$$\mathbf{w} = \operatorname{argmax}_{\|\mathbf{w}\|=1} (\mathbf{X}\mathbf{w})^T (\mathbf{X}\mathbf{w})$$

$$= \operatorname{argmax}_{\|\mathbf{w}\|=1} \mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w}$$

$$= \operatorname{argmax}_{\|\mathbf{w}\|=1} \mathbf{w}^T S_{XX} \mathbf{w}$$

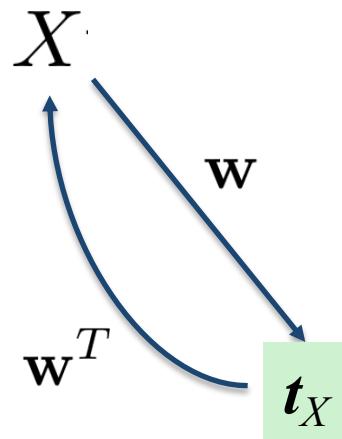


$$\mathcal{L}(\mathbf{w}, \lambda) = \mathbf{w}^T S_{XX} \mathbf{w} - \lambda [\mathbf{w}_X^T \mathbf{w}_X - 1]$$

$$S_{XX} \mathbf{w} = \lambda \mathbf{w}$$

Non-linear iterative partial least squares - NIPALS

Wold 1975



1. Random initialization \mathbf{w}
2. Solve :
$$\operatorname{argmin}_{\mathbf{t}} \|\mathbf{X} - \mathbf{t}\mathbf{w}^T\|^2$$
$$\mathbf{t} = \mathbf{X}\mathbf{w}(\mathbf{w}^T\mathbf{w})^{-1}$$
3. Normalize :
$$\mathbf{t} = \frac{\mathbf{t}}{\|\mathbf{t}\|}$$
4. Update :
$$\operatorname{argmin}_{\mathbf{w}} \|\mathbf{X} - \mathbf{t}\mathbf{w}^T\|^2$$
$$\mathbf{w} = \mathbf{X}^T\mathbf{t}(\mathbf{t}^T\mathbf{t})^{-1}$$
5. Iterate 2-4 until convergence

Why it works:

$$4 \rightarrow \text{const } \mathbf{w} = \mathbf{X}^T\mathbf{t}$$

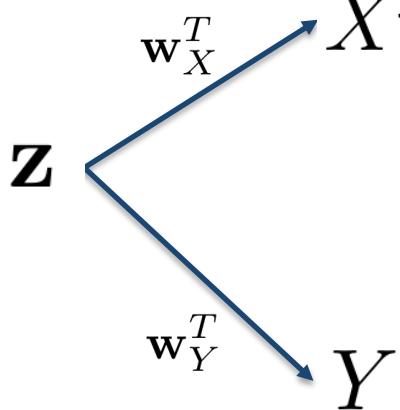
$$2 \rightarrow \text{const } \mathbf{t} = \mathbf{X}\mathbf{w}$$

Then

$$\text{const } \mathbf{w} = S_{XX}\mathbf{w}$$

eigen-solution of the covariance matrix

Multi-modal latent variable models



The equation block illustrates the decomposition of the matrices X and Y into latent variable Z , weight vectors w_X^T and w_Y^T , and error terms ϵ_X and ϵ_Y .

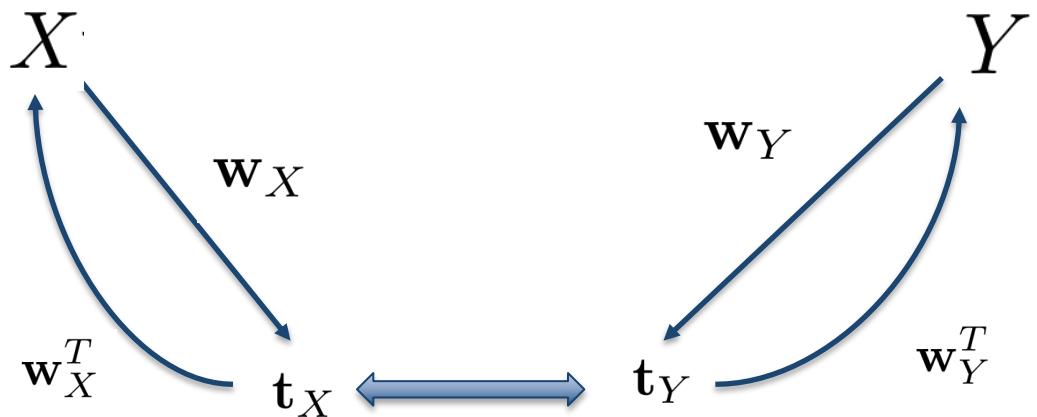
$$X = \begin{matrix} Z \\ \vdots \end{matrix} \begin{matrix} w_X^T \\ + \end{matrix} \epsilon_X$$

inference

$$Y = \begin{matrix} Z \\ \vdots \end{matrix} \begin{matrix} w_Y^T \\ + \end{matrix} \epsilon_Y$$

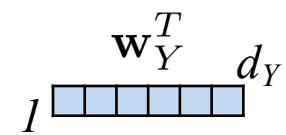
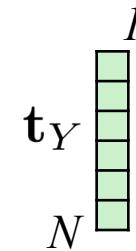
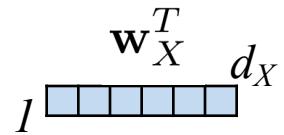
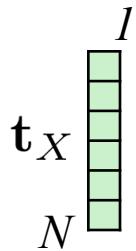
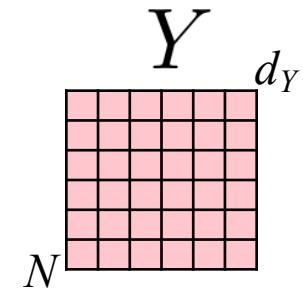
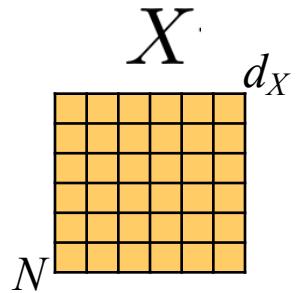
Multi-modal latent variable models

$$X = \begin{matrix} t_X \\ \vdots \\ t_X \end{matrix} \begin{matrix} w_X^T \\ \vdots \\ w_X^T \end{matrix} + \varepsilon_X$$
$$Y = \begin{matrix} t_Y \\ \vdots \\ t_Y \end{matrix} \begin{matrix} w_Y^T \\ \vdots \\ w_Y^T \end{matrix} + \varepsilon_Y$$

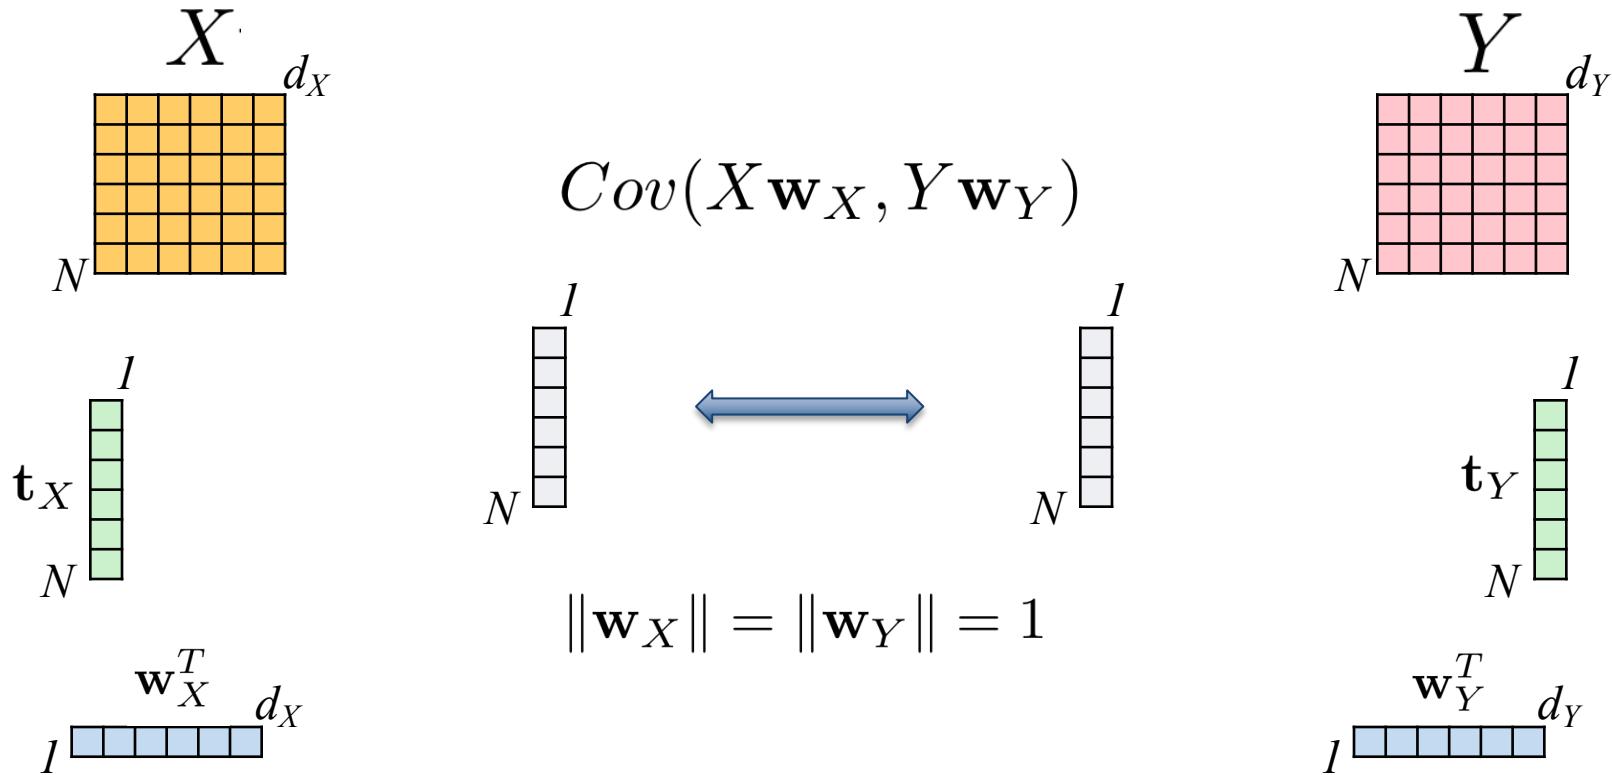


Goal:
Identifying
- Projections
(loadings) \mathbf{W}
- Latent
representation
(scores) \mathbf{t}

Multi-modal latent variable models



Multi-modal latent variable models



Partial Least Squares

A **covariance** maximisation problem:

$$\operatorname{argmax}_{\mathbf{w}_X, \mathbf{w}_Y} \operatorname{Cov}(X\mathbf{w}_X, Y\mathbf{w}_Y)$$

$$\operatorname{Cov}(X\mathbf{w}_X, Y\mathbf{w}_Y) = \frac{\mathbf{w}_X^T S_{XY} \mathbf{w}_Y}{\sqrt{\mathbf{w}_X^T \mathbf{w}_X} \sqrt{\mathbf{w}_Y^T \mathbf{w}_Y}}$$

Partial Least Squares

$$\mathcal{L}(\mathbf{w}_X, \mathbf{w}_Y, \lambda_X, \lambda_Y) = \mathbf{w}_X^T S_{XY} \mathbf{w}_Y - \lambda_X [\mathbf{w}_X^T \mathbf{w}_X - 1] - \lambda_Y [\mathbf{w}_Y^T \mathbf{w}_Y - 1]$$

$$\begin{cases} S_{XY} \mathbf{w}_Y = \lambda_X \mathbf{w}_X \\ S_{YX} \mathbf{w}_X = \lambda_Y \mathbf{w}_Y \end{cases} \quad \xrightarrow{\hspace{1cm}} \quad \begin{aligned} \lambda_X \mathbf{w}_X^T \mathbf{w}_X &= \mathbf{w}_X^T S_{XY} \mathbf{w}_Y \\ &= \mathbf{w}_Y^T S_{YX} \mathbf{w}_X \\ &= \lambda_Y \mathbf{w}_Y^T \mathbf{w}_Y \end{aligned}$$

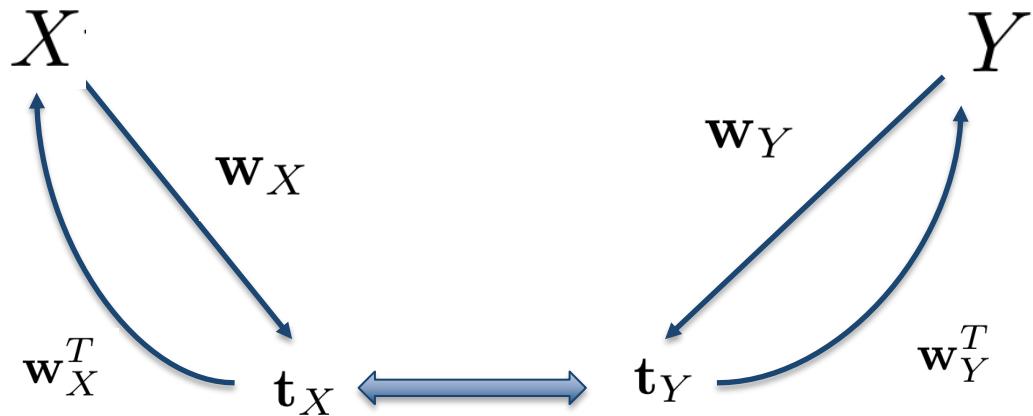
$$\lambda_X = \lambda_Y = \lambda$$

$$\begin{bmatrix} 0 & S_{XY} \\ S_{YX} & 0 \end{bmatrix} \begin{bmatrix} \mathbf{w}_X \\ \mathbf{w}_Y \end{bmatrix} = \lambda \begin{bmatrix} \mathbf{w}_X \\ \mathbf{w}_Y \end{bmatrix}$$

The PLS problem is solved via singular value decomposition (SVD) of the covariance matrix

Non-linear iterative partial least squares - NIPALS

[scikit-learn/sklearn/cross_decomposition](#)



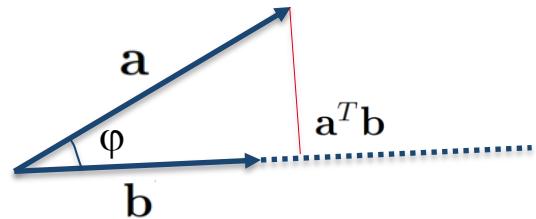
1. Random initialization \mathbf{t}_X
2. Update \mathbf{w}_Y :
$$\operatorname{argmin}_{\mathbf{w}_Y} \| \mathbf{Y} - \mathbf{t}_X \mathbf{w}_Y^T \|^2$$
$$\mathbf{w}_Y = \mathbf{Y}^T \mathbf{t}_X (\mathbf{t}_X^T \mathbf{t}_X)^{-1}$$
3. Normalize :
$$\mathbf{w}_Y = \frac{\mathbf{w}_Y}{\|\mathbf{w}_Y\|}$$
4. $\mathbf{t}_Y = \mathbf{Y} \mathbf{w}_Y$
5. Update \mathbf{w}_X :
$$\operatorname{argmin}_{\mathbf{w}_X} \| \mathbf{X} - \mathbf{t}_Y \mathbf{w}_X^T \|^2$$
$$\mathbf{w}_X = \mathbf{X}^T \mathbf{t}_Y (\mathbf{t}_Y^T \mathbf{t}_Y)^{-1}$$
6. Normalize :
$$\mathbf{w}_X = \frac{\mathbf{w}_X}{\|\mathbf{w}_X\|}$$
7. $\mathbf{t}_X = \mathbf{X} \mathbf{w}_X$
8. Iterate 2-7 until convergence

Canonical Correlation Analysis

A **correlation** maximisation problem:

$$\operatorname{argmax}_{\mathbf{w}_X, \mathbf{w}_Y} \rho(X\mathbf{w}_X, Y\mathbf{w}_Y)$$

$$\rho(\mathbf{a}, \mathbf{b}) = \mathbf{a}^T \mathbf{b} / (\sqrt{\mathbf{a}^T \mathbf{a}} \sqrt{\mathbf{b}^T \mathbf{b}})$$



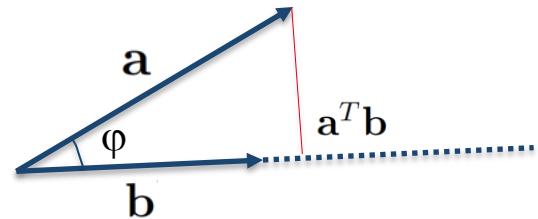
$$\rho(X\mathbf{w}_X, Y\mathbf{w}_Y) = \frac{\mathbf{w}_X^T S_{XY} \mathbf{w}_Y}{\sqrt{\mathbf{w}_X^T S_{XX} \mathbf{w}_X} \sqrt{\mathbf{w}_Y^T S_{YY} \mathbf{w}_Y}}$$

Canonical Correlation Analysis

A **correlation** maximisation problem:

$$\operatorname{argmax}_{\mathbf{w}_X, \mathbf{w}_Y} \rho(X\mathbf{w}_X, Y\mathbf{w}_Y)$$

$$\rho(\mathbf{a}, \mathbf{b}) = \mathbf{a}^T \mathbf{b} / (\sqrt{\mathbf{a}^T \mathbf{a}} \sqrt{\mathbf{b}^T \mathbf{b}})$$



$$\rho(X\mathbf{w}_X, Y\mathbf{w}_Y) = \frac{\mathbf{w}_X^T S_{XY} \mathbf{w}_Y}{\sqrt{\mathbf{w}_X^T S_{XX} \mathbf{w}_X} \sqrt{\mathbf{w}_Y^T S_{YY} \mathbf{w}_Y}}$$

Cross-covariance

Within-modality covariance

Canonical Correlation Analysis

$$\mathcal{L}(\mathbf{w}_X, \mathbf{w}_Y, \lambda_X, \lambda_Y) = \mathbf{w}_X^T S_{XY} \mathbf{w}_Y - \lambda_X [\mathbf{w}_X^T S_{XX} \mathbf{w}_X - 1] - \lambda_Y [\mathbf{w}_Y^T S_{YY} \mathbf{w}_Y - 1]$$

Canonical Correlation Analysis

$$\mathcal{L}(\mathbf{w}_X, \mathbf{w}_Y, \lambda_X, \lambda_Y) = \mathbf{w}_X^T S_{XY} \mathbf{w}_Y - \lambda_X [\mathbf{w}_X^T S_{XX} \mathbf{w}_X - 1] - \lambda_Y [\mathbf{w}_Y^T S_{YY} \mathbf{w}_Y - 1]$$

$$\begin{cases} S_{XY} \mathbf{w}_Y = \lambda_X S_{XX} \mathbf{w}_X \\ S_{YX} \mathbf{w}_X = \lambda_Y S_{YY} \mathbf{w}_Y \end{cases}$$

Canonical Correlation Analysis

$$\mathcal{L}(\mathbf{w}_X, \mathbf{w}_Y, \lambda_X, \lambda_Y) = \mathbf{w}_X^T S_{XY} \mathbf{w}_Y - \lambda_X [\mathbf{w}_X^T S_{XX} \mathbf{w}_X - 1] - \lambda_Y [\mathbf{w}_Y^T S_{YY} \mathbf{w}_Y - 1]$$

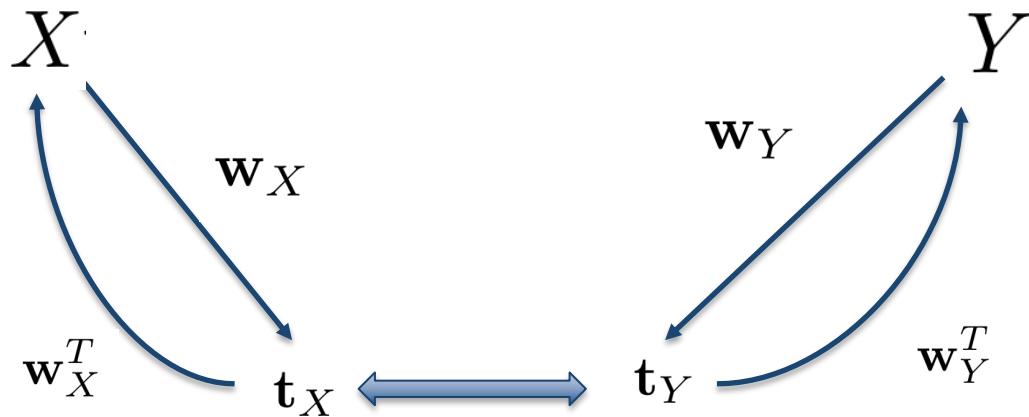
$$\begin{cases} S_{XY} \mathbf{w}_Y = \lambda_X S_{XX} \mathbf{w}_X \\ S_{YX} \mathbf{w}_X = \lambda_Y S_{YY} \mathbf{w}_Y \end{cases}$$

$$\begin{bmatrix} 0 & S_{XY} \\ S_{YX} & 0 \end{bmatrix} \begin{bmatrix} \mathbf{w}_X \\ \mathbf{w}_Y \end{bmatrix} = \lambda \begin{bmatrix} S_{XX} & 0 \\ 0 & S_{YY} \end{bmatrix} \begin{bmatrix} \mathbf{w}_X \\ \mathbf{w}_Y \end{bmatrix}$$

CCA is solved as a
generalized
eigenvalue problem

Non-linear iterative partial least squares - NIPALS

[scikit-learn](#)/[sklearn](#)/[cross_decomposition](#)



1. Random initialization \mathbf{t}_X
2. Update \mathbf{w}_Y :
$$\operatorname{argmin}_{\mathbf{w}_Y} \|\mathbf{Y}\mathbf{w}_Y - \mathbf{t}_X\|^2$$
$$\mathbf{w}_Y = (\mathbf{Y}^T \mathbf{Y})^{-1} \mathbf{Y}^T \mathbf{t}_X$$
3. Normalize :
$$\mathbf{w}_Y = \frac{\mathbf{w}_Y}{\|\mathbf{w}_Y\|}$$
4. $\mathbf{t}_Y = \mathbf{Y}\mathbf{w}_Y$
5. Update \mathbf{w}_X :
$$\operatorname{argmin}_{\mathbf{w}_X} \|\mathbf{X}\mathbf{w}_X - \mathbf{t}_Y\|^2$$
$$\mathbf{w}_X = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{t}_Y$$
6. Normalize :
$$\mathbf{w}_X = \frac{\mathbf{w}_X}{\|\mathbf{w}_X\|}$$
7. $\mathbf{t}_X = \mathbf{X}\mathbf{w}_X$
8. Iterate 2-7 until convergence

Non-linear iterative partial least squares - NIPALS Deflation

$$\mathbf{X}^{(i+1)} = \mathbf{X}^{(i)} - \mathbf{t}^{(i)} \frac{\mathbf{t}^{(i)T} \mathbf{X}^{(i)}}{\mathbf{t}^{(i)T} \mathbf{t}^{(i)}},$$
$$\mathbf{Y}^{(i+1)} = \mathbf{Y}^{(i)} - \mathbf{u}^{(i)} \frac{\mathbf{u}^{(i)T} \mathbf{Y}^{(i)}}{\mathbf{u}^{(i)T} \mathbf{u}^{(i)}}.$$

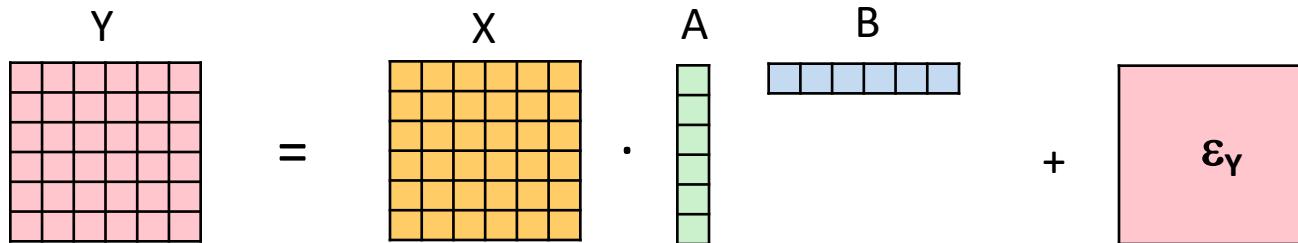
Iterate until

- residual component negligible epsilon
- Difference between consecutive residual components negligible

Reduced Rank Regression

$$\mathbf{Y} = \mathbf{X}\mathbf{C} + \epsilon.$$

$$\mathbf{C} = \mathbf{A}\mathbf{B}$$

$$\mathbf{Y} = \mathbf{X} \cdot \mathbf{A} \mathbf{B} + \epsilon_Y$$


$$f(\mathbf{A}, \mathbf{B}) = \text{tr}\{(\mathbf{Y} - \mathbf{XAB})\Gamma(\mathbf{Y} - \mathbf{XAB})^T\}$$

Reduced Rank Regression

Solution associated to the eigen-decomposition of the matrix

$$\mathbf{R} = \boldsymbol{\Gamma}^{1/2} \mathbf{S}_{\mathbf{YX}} \mathbf{S}_{\mathbf{XX}}^{-1} \mathbf{S}_{\mathbf{XY}} \boldsymbol{\Gamma}^{1/2}$$

Matrix encoding
prior knowledge
on \mathbf{Y}

Reduced Rank Regression

Solution associated to the eigen-decomposition of the matrix

$$\mathbf{R} = \boldsymbol{\Gamma}^{1/2} \mathbf{S}_{\mathbf{YX}} \mathbf{S}_{\mathbf{XX}}^{-1} \mathbf{S}_{\mathbf{XY}} \boldsymbol{\Gamma}^{1/2}$$

Matrix encoding prior knowledge on \mathbf{Y}

RRR solutions:

$$\mathbf{A} = \boldsymbol{\Gamma}^{-1/2} \mathbf{U}, \quad \mathbf{B} = \mathbf{U}^T \boldsymbol{\Gamma}^{1/2} \mathbf{S}_{\mathbf{YX}} \mathbf{S}_{\mathbf{XX}}^{-1}$$

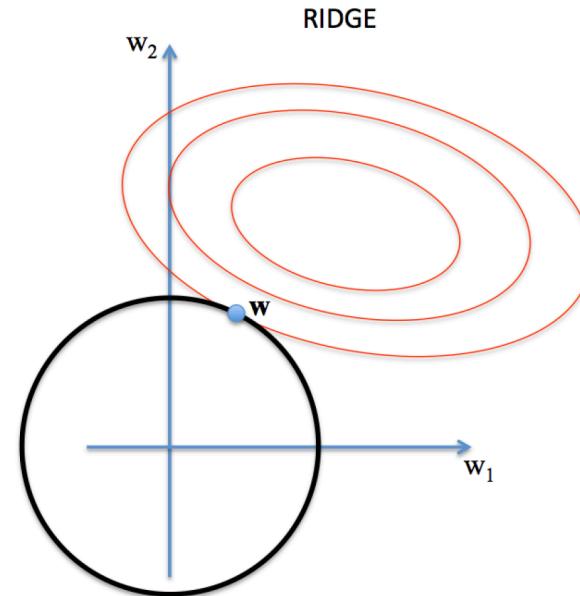
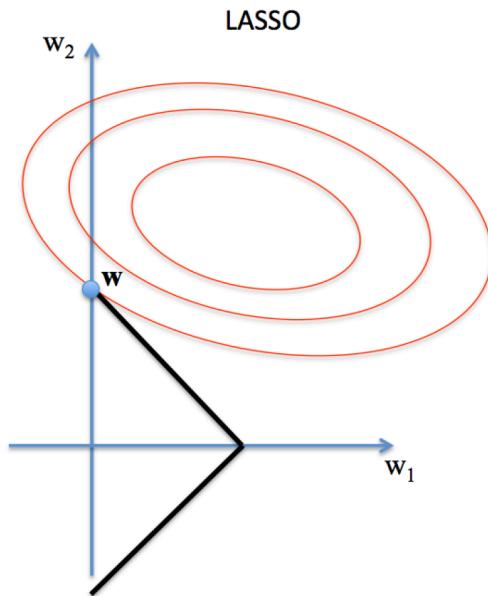
Special case:

$$\boldsymbol{\Gamma} = \mathbf{S}_{\mathbf{YY}} \quad \xrightarrow{\hspace{1cm}} \quad \text{CCA}$$

Sparsity in latent variable models

$$\underset{\mathbf{w}}{\operatorname{argmin}} f(\mathbf{w}) + \lambda \|\mathbf{w}\|_1$$

$$\underset{\mathbf{w}}{\operatorname{argmin}} f(\mathbf{w}) + \lambda \|\mathbf{w}\|_2^2$$



Algorithm Regularization of projections parameters \mathbf{w}_x and \mathbf{w}_y in NIPALS.

Given current estimates of \mathbf{w}_x and \mathbf{w}_y .

While not converged do:

1. compute $\mathbf{t} = \mathbf{X}\mathbf{w}_x$,

2. compute $\mathbf{u} = \mathbf{Y}\mathbf{w}_y$,

3. compute $\overline{\mathbf{w}}_x$ by solving the Elastic-Net regression:

$$\overline{\mathbf{w}}_x = \arg \min_{\mathbf{v}} (\mathbf{t} - \mathbf{X}\mathbf{v})^2 + \lambda_{x2}\|\mathbf{v}\|_2^2 + \lambda_{x1}\|\mathbf{v}\|_1,$$

4. compute $\overline{\mathbf{w}}_y$ by solving the Elastic-Net regression:

$$\overline{\mathbf{w}}_y = \arg \min_{\mathbf{v}} (\mathbf{u} - \mathbf{Y}\mathbf{v})^2 + \lambda_{y2}\|\mathbf{v}\|_2^2 + \lambda_{y1}\|\mathbf{v}\|_1,$$

3. Normalize $\overline{\mathbf{w}}_x$ and $\overline{\mathbf{w}}_y$,

4. Set $\mathbf{w}_x = \overline{\mathbf{w}}_x$, $\mathbf{w}_y = \overline{\mathbf{w}}_y$.

PLS in practice

Application to imaging-genetics analysis
in Alzheimer's disease

Multivariate Association studies

Maximizing the joint relationship between genetic variants and brain features

$$\mathbf{X} = \begin{matrix} \text{~10}^6 \text{ SNPs} \\ \text{N individuals} \end{matrix}$$
$$\mathbf{Y} = \begin{matrix} \text{~10}^5 \text{ brain features} \\ \text{N individuals} \end{matrix}$$

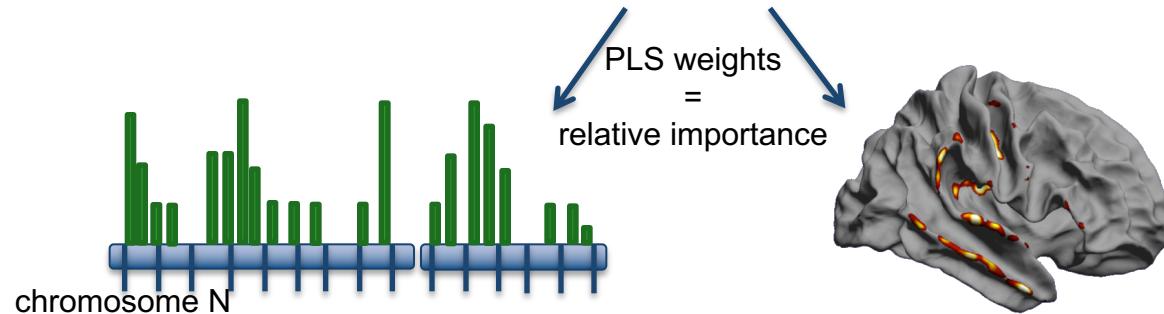
Partial least squares (PLS)
 $\max_{\mathbf{p}, \mathbf{q}} \text{Cov}(\mathbf{X} \cdot \mathbf{p}, \mathbf{Y} \cdot \mathbf{q})$

Multivariate Association studies

Maximizing the joint relationship between genetic variants and brain features

$$X = \begin{matrix} \text{~10}^6 \text{ SNPs} \\ \text{---} \\ N \text{ individuals} \end{matrix}$$
$$Y = \begin{matrix} \text{~10}^5 \text{ brain features} \\ \text{---} \\ N \text{ individuals} \end{matrix}$$

Partial least squares (PLS)
 $\max_{p,q} \text{Cov}(X \cdot p, Y \cdot q)$

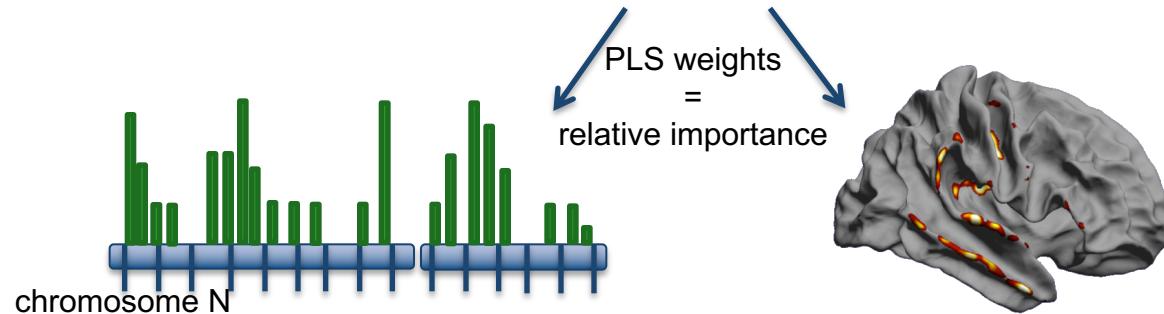


Multivariate Association studies

Maximizing the joint relationship between genetic variants and brain features

$$X = \begin{matrix} \text{~10}^6 \text{ SNPs} \\ \text{N individuals} \end{matrix}$$
$$Y = \begin{matrix} \text{~10}^5 \text{ brain features} \\ \text{N individuals} \end{matrix}$$

Partial least squares (PLS)
 $\max_{p,q} \text{Cov}(X \cdot p, Y \cdot q)$



- Pros.** Overcomes issues of mass univariate analysis
- Avoiding independent **multiple testing**
 - Exploring **SNP-SNP interaction** (epistatic effects)

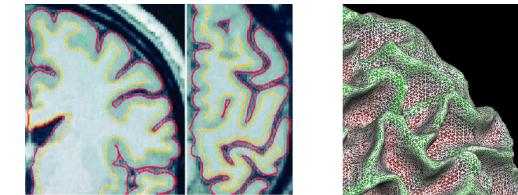
Study cohort

	Healthy	AD
N	401	238
Age (years)	74.45	74.72
Sex (% females)	49	45
MMSE	29.1	23.2
Apoe4 (% 0/1/2)	72/26/2	31/48/21



X = Phenotype features

- Freesurfer **brain cortical thickness** maps (327,684 mesh points)
- **Radial distance** of hippocampi and amygdalae (27,120 mesh points) [Gutman et al, NeuroImage 2013]



Y = Genotype features

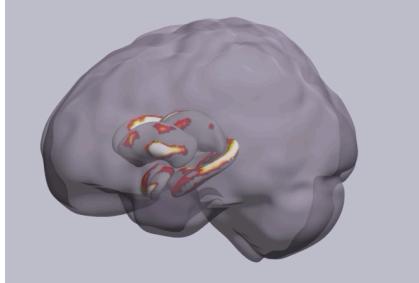
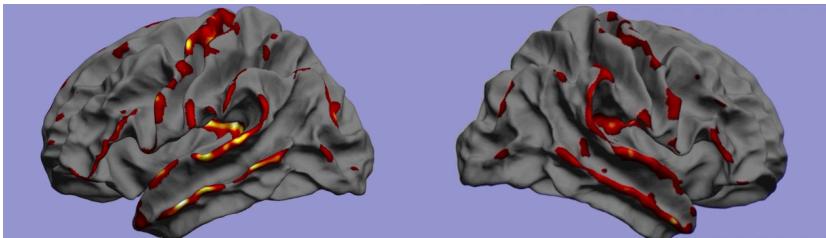
- Individuals' minor allele counts for **1,167,126 SNPs** in chromosomes 1 to 22

Standard quality control: MAF < 0.01, Genotype Call Rate <95% , Hardy-Weinberg Equilibrium < 1×10^{-6} .

Imputation to HapMap III reference panel, quality controlled (MAF > 0.01 and R-squared > 0.3)

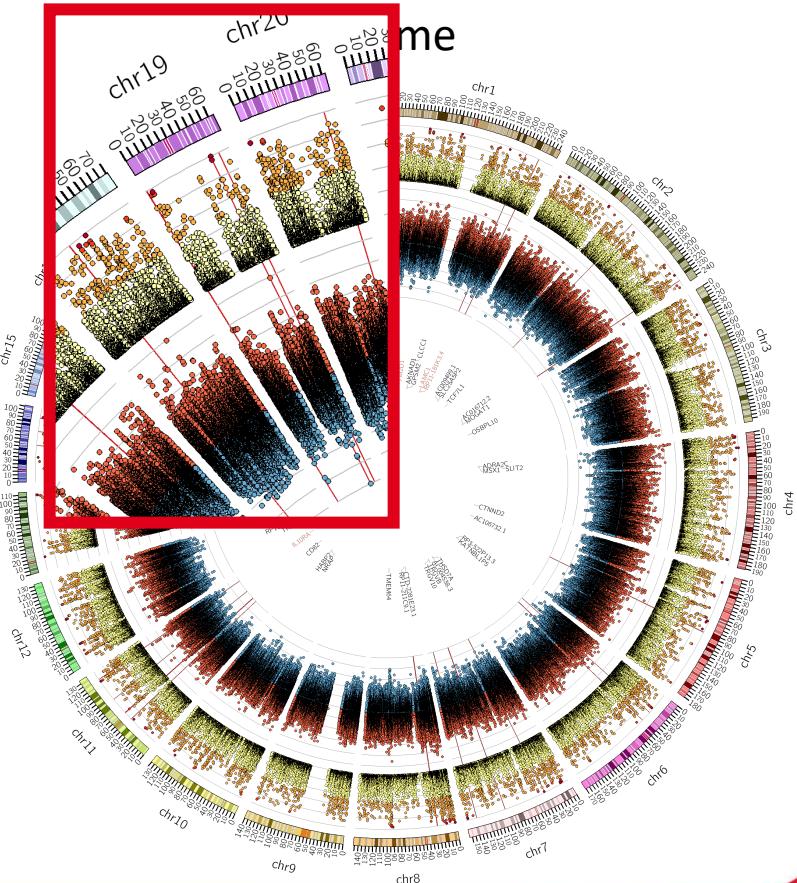
Application to multivariate Imaging-genetics

Atrophy profile from brain imaging



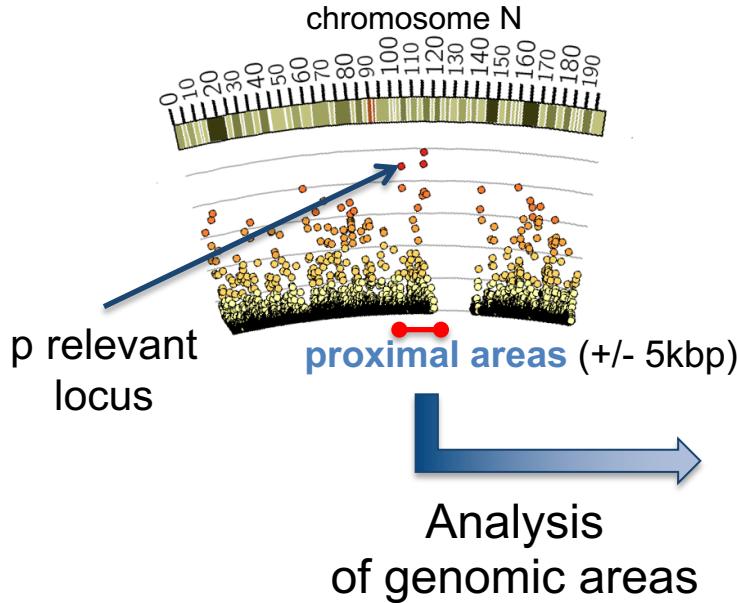
↔
Partial least
squares
association

639 individuals
401 healthy
238 Alzheimer's



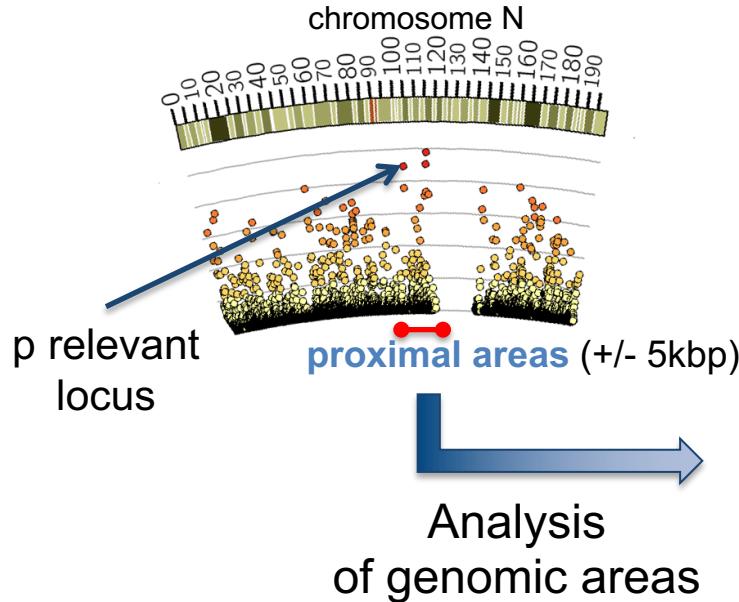
Investigating biological mechanisms through Meta-analysis

PLS statistical result



Investigating biological mechanisms through Meta-analysis

PLS statistical result



Querying gene annotation databases



McLaren et al. The Ensembl Variant Effect Predictor. Genome Biology, 2016

Investigating biological mechanisms through Meta-analysis



148 SNP-gene combinations

6 tested tissues

*hippocampus, whole blood,
Adipose subcutaneous, artery tibia, nerve tibial,
treated fibroblast*

14 Significantly expressed genes

TM2D1 (amyloid-beta binding protein),
IL10RA (increase in hippo in mouse model),
TRIB3
(neuronal cell death, modulates PSEN1 stability, interacts with APP)

	Significance (p-value) training	Significance (p-value) testing
TM2D1	0.005	0.053
IL10RA	0.107	0.620
TRIB3	0.003	0.003
ZBTB7A	0.036	0.913
LYSMD4	0.000	0.206
CRYL1	0.621	0.118
FAM135B	0.000	0.559
IP6K3	0.000	0.465
ITGA1	0.099	0.731
KIN	0.001	0.206
LAMC1	0.002	0.062
LINC00941	0.000	0.690
RBPMS2	0.000	0.215
RP11-181K3.4	0.002	0.053

Latent variable models

1. *Multi-variate modeling*
2. Novel scalable approaches to *multi-view* data

Latent variable models via Variational Autoencoders

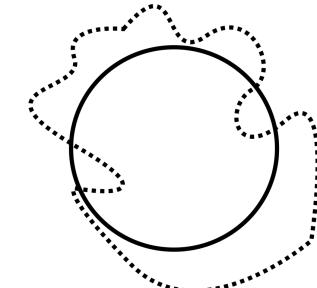
Kingma & Welling, 2014; Rezende et al. 2014

$$\begin{array}{ccc} \mathbf{z} & \xrightarrow{\hspace{2cm}} & \mathbf{x} \\ \text{Posterior} & p(\mathbf{z}|\mathbf{x}) & p(\mathbf{x}|\mathbf{z}) \quad \text{Likelihood} \end{array}$$

$$p(\mathbf{z}|\mathbf{x}) = \int_{\mathbf{z}} p(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z}$$

Difficult to compute

..... $p(\mathbf{z}|\mathbf{x})$
— $q(\mathbf{z}|\mathbf{x})$



Idea: find a “close enough” and simple approximation $q(\mathbf{z}|\mathbf{x})$

Latent variable models via Variational Autoencoders

Kingma & Welling, 2014; Rezende et al. 2014

$$\mathbf{z} \longrightarrow \mathbf{x}$$

$$D_{KL}[q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})] = \mathbb{E}_{\mathbf{z} \sim q} \log[q(\mathbf{z}|x)] - \mathbb{E}_{\mathbf{z} \sim q} \log[p(\mathbf{z}|x)]$$

Latent variable models via Variational Autoencoders

Kingma & Welling, 2014; Rezende et al. 2014

Z → **X**

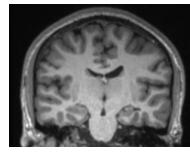
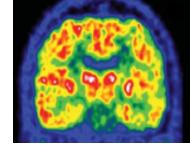
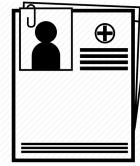
$$D_{KL}[q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})] = \mathbf{E}_{\mathbf{z} \sim q} \log[q(\mathbf{z}|x)] - \mathbf{E}_{\mathbf{z} \sim q} \log[p(\mathbf{z}|x)]$$

Evidence lower bound (ELBO)

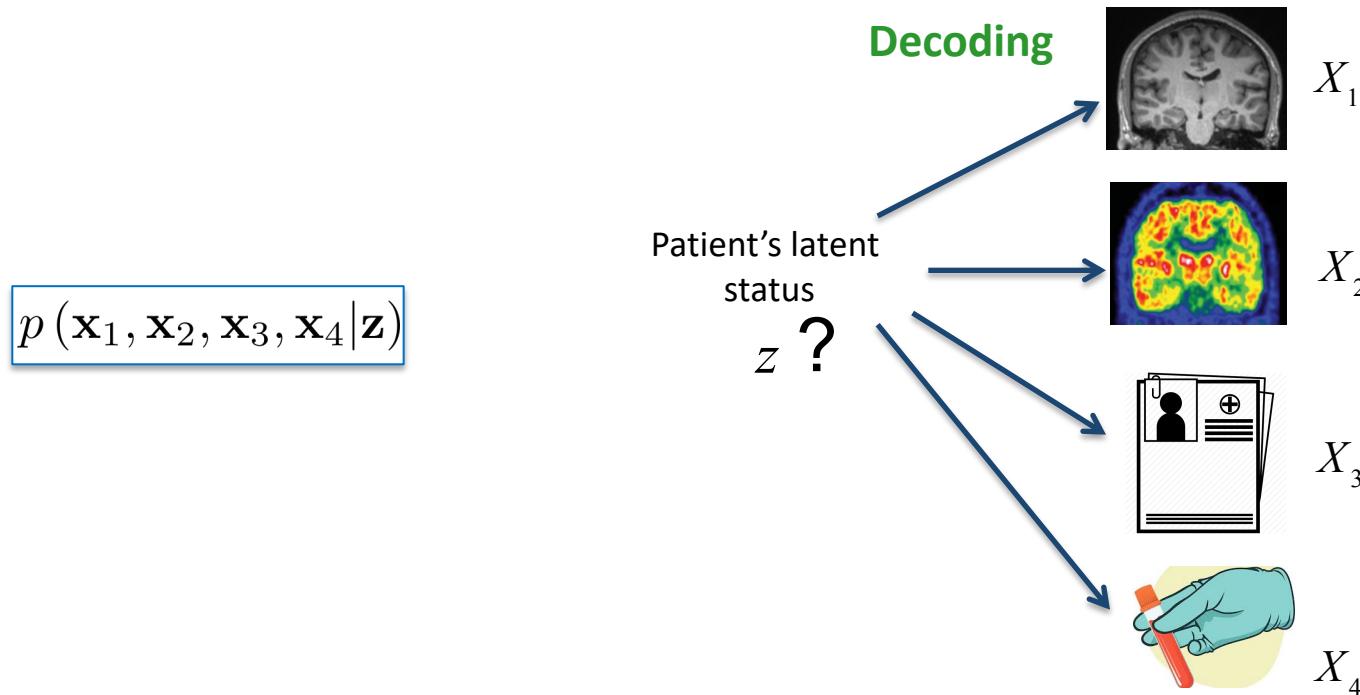
$$\mathcal{L} = \mathbf{E}_{\mathbf{z} \sim q} \log[p(\mathbf{x}|\mathbf{z})] - D_{KL}[q(\mathbf{z}|\mathbf{x})||p(\mathbf{z})]$$

reconstruction regularization

Generative representation of multimodal data

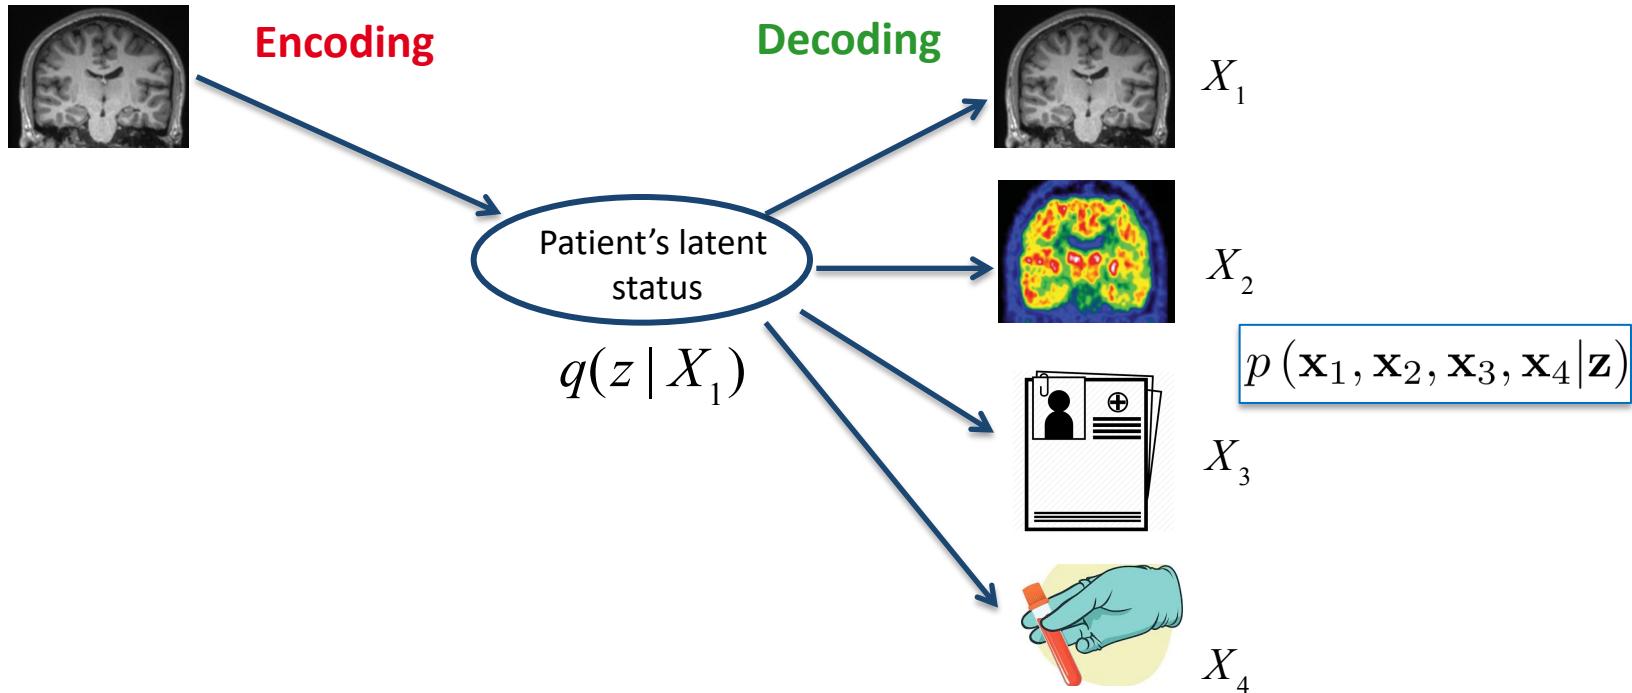
 X_1  X_2  X_3  X_4

Generative representation of multimodal data



Decoding: data reconstruction from the latent representation

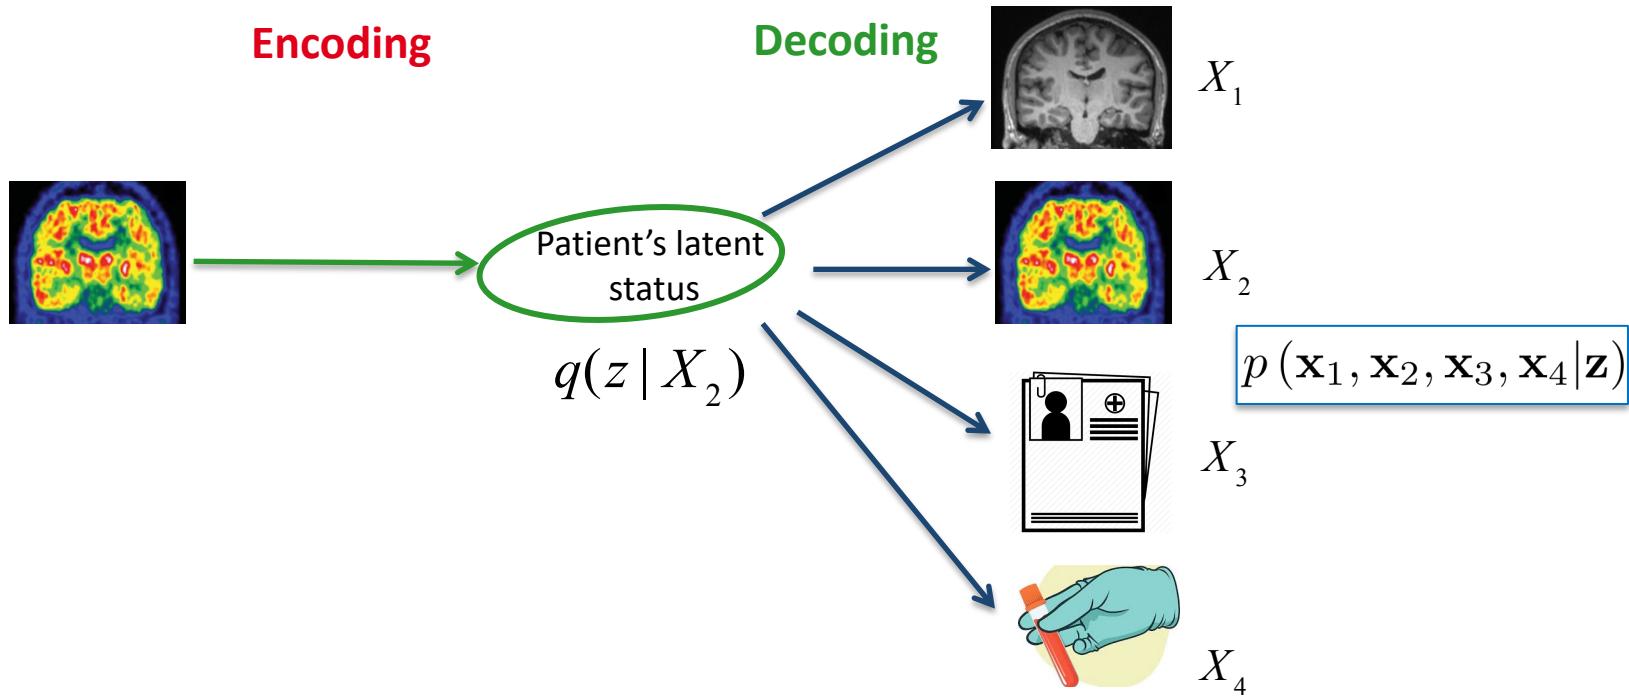
Generative representation of multimodal data



Decoding: data reconstruction from the latent representation

Encoding: latent representation from the data

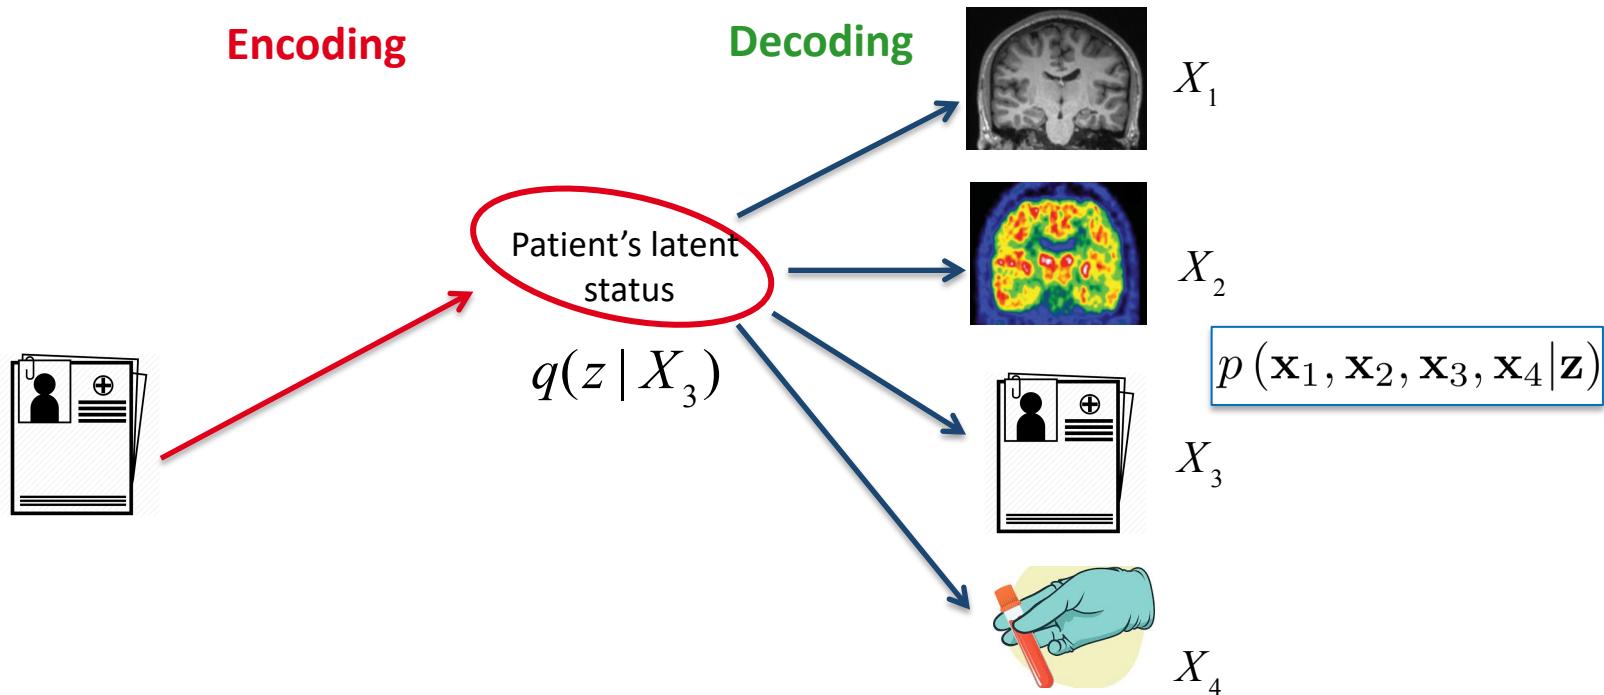
Generative representation of multimodal data



Decoding: data reconstruction from the latent representation

Encoding: latent representation from the data

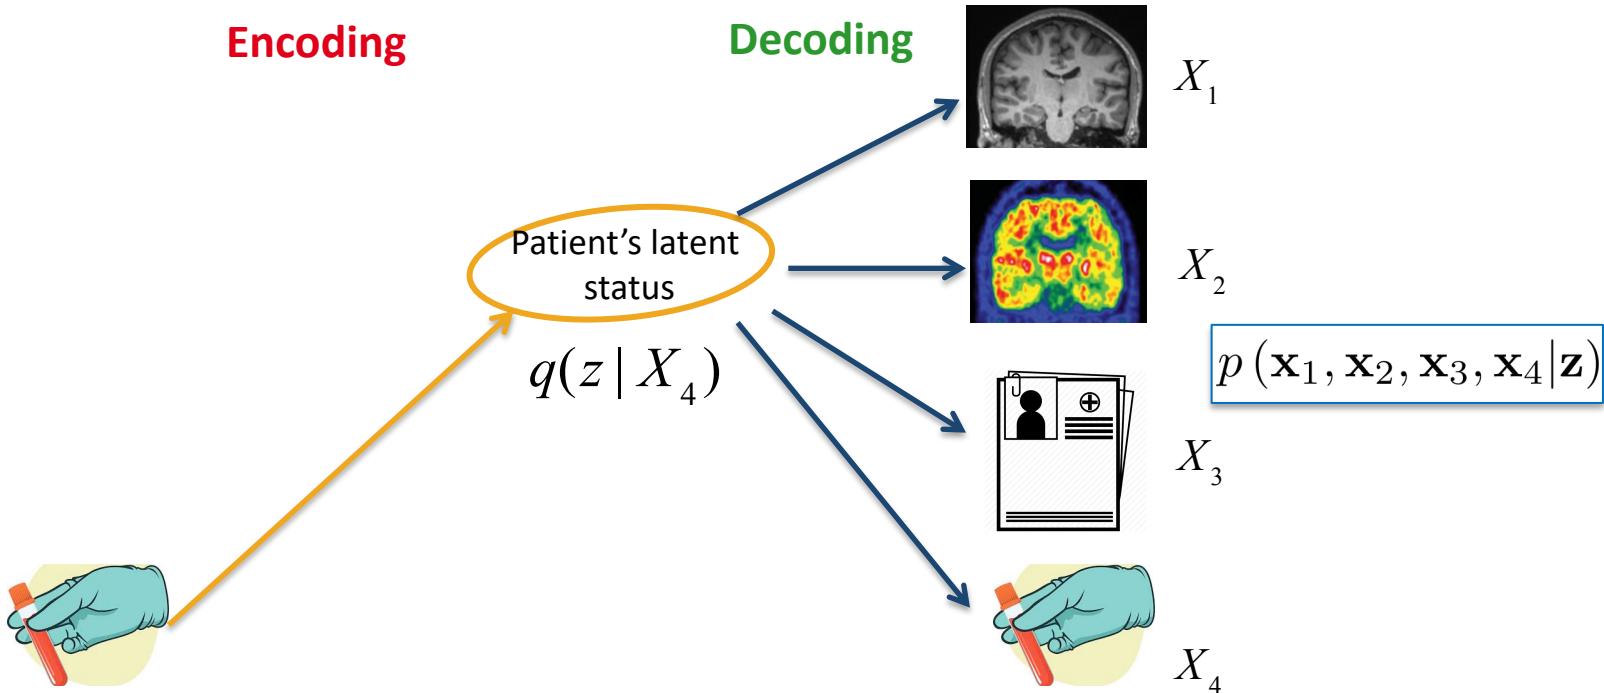
Generative representation of multimodal data



Decoding: data reconstruction from the latent representation

Encoding: latent representation from the data

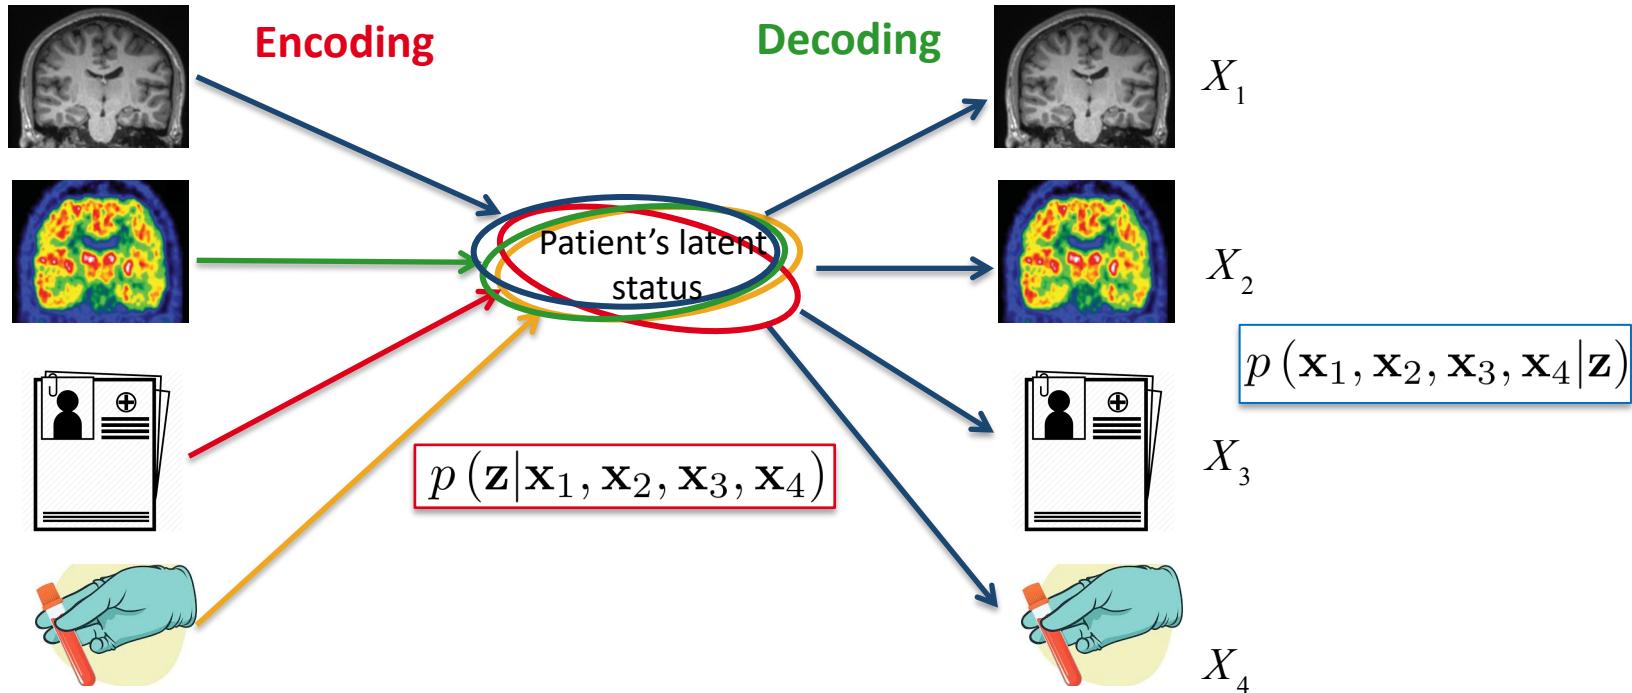
Generative representation of multimodal data



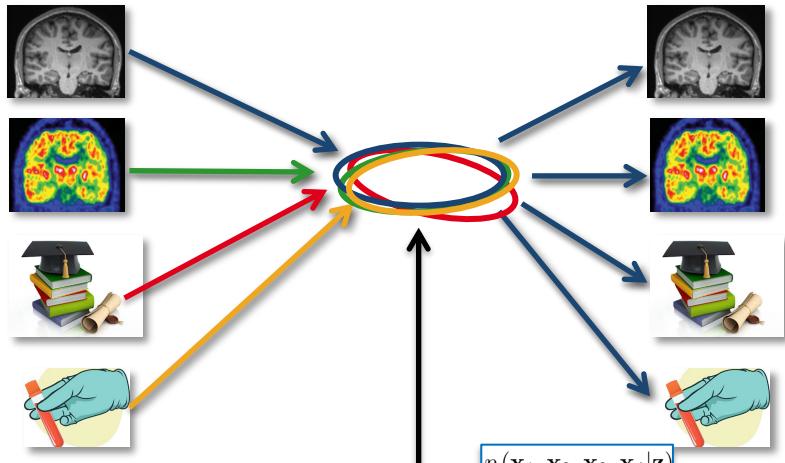
Decoding: data reconstruction from the latent representation

Encoding: latent representation from the data

Generative representation of multimodal data



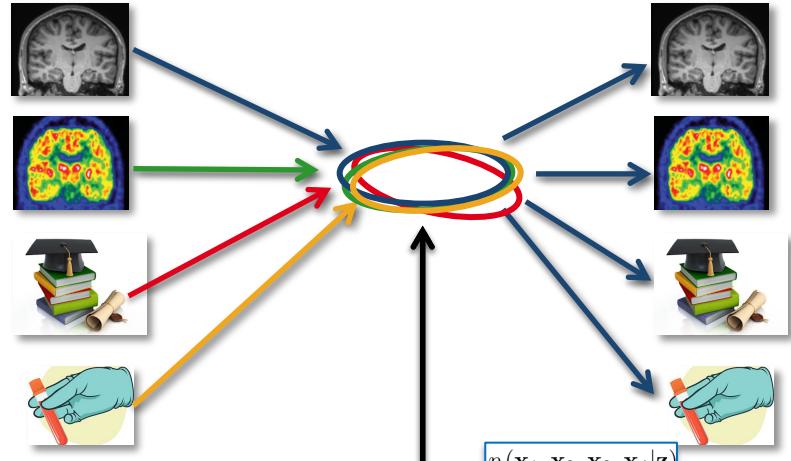
Generative representation of multimodal data



minimize

$$\frac{1}{C} \mathcal{D}_{\text{KL}} \left(q(\mathbf{z} | \mathbf{x}_c) || p(\mathbf{z} | \mathbf{x}_1, \dots, \mathbf{x}_C) \right)$$

Generative representation of multimodal data



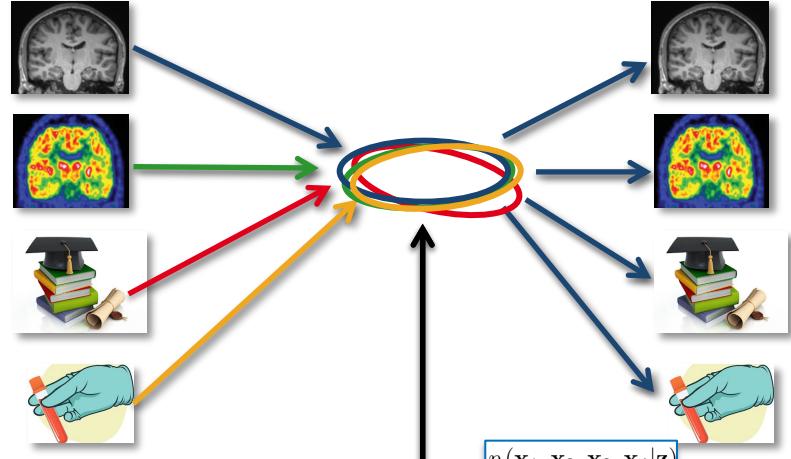
Evidence Lower bound (ELBO)

$$\frac{1}{C} \sum_{c=1}^C \mathbb{E}_{q(\mathbf{z}|\mathbf{x}_c)} \left[\sum_{i=1}^C \ln p(\mathbf{x}_i|\mathbf{z}) \right] - \mathcal{D}_{\text{KL}}(q(\mathbf{z}|\mathbf{x}_c) || p(\mathbf{z}))$$

minimize

$$\frac{1}{C} \mathcal{D}_{\text{KL}}(q(\mathbf{z}|\mathbf{x}_c) || p(\mathbf{z}|\mathbf{x}_1, \dots, \mathbf{x}_C))$$

Generative representation of multimodal data



Evidence Lower bound (ELBO)

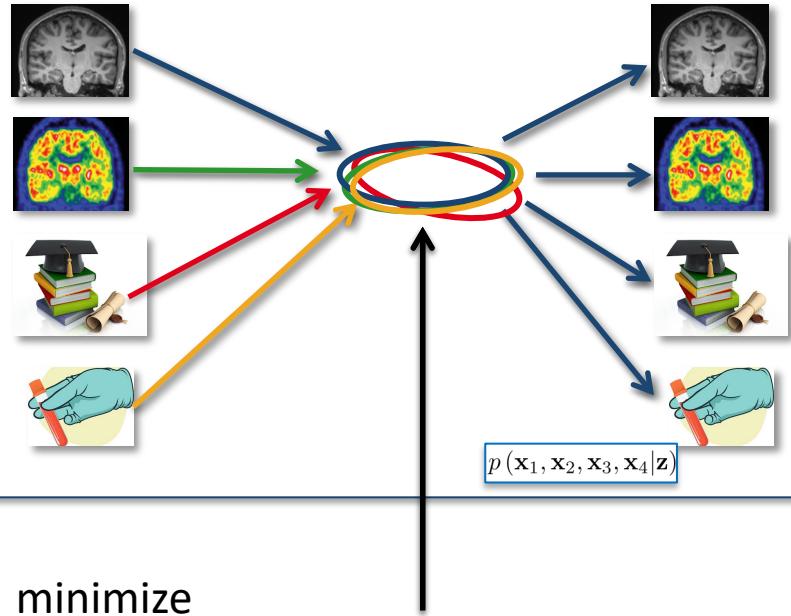
$$\frac{1}{C} \sum_{c=1}^C \mathbb{E}_{q(\mathbf{z}|\mathbf{x}_c)} \left[\sum_{i=1}^C \ln p(\mathbf{x}_i | \mathbf{z}) \right] - \mathcal{D}_{\text{KL}}(q(\mathbf{z}|\mathbf{x}_c) || p(\mathbf{z}))$$

Encoding for given channel

minimize

$$\frac{1}{C} \mathcal{D}_{\text{KL}}(q(\mathbf{z}|\mathbf{x}_c) || p(\mathbf{z}|\mathbf{x}_1, \dots, \mathbf{x}_C))$$

Generative representation of multimodal data



minimize

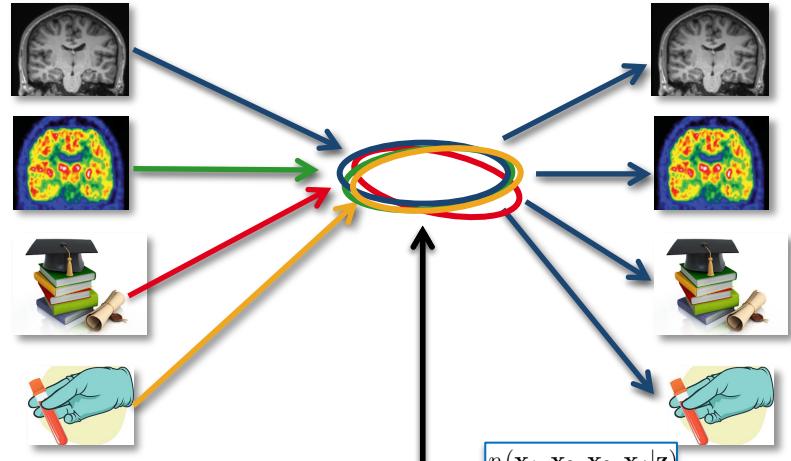
$$\frac{1}{C} \mathcal{D}_{\text{KL}}(q(\mathbf{z}|\mathbf{x}_c) || p(\mathbf{z}|\mathbf{x}_1, \dots, \mathbf{x}_C))$$

Evidence Lower bound (ELBO)

$$\frac{1}{C} \sum_{c=1}^C \mathbb{E}_{q(\mathbf{z}|\mathbf{x}_c)} \left[\sum_{i=1}^C \ln p(\mathbf{x}_i | \mathbf{z}) \right] - \mathcal{D}_{\text{KL}}(q(\mathbf{z}|\mathbf{x}_c) || p(\mathbf{z}))$$

Encoding for given channel
Reconstruction of all channels

Generative representation of multimodal data



minimize

$$\frac{1}{C} \mathcal{D}_{\text{KL}} (q(\mathbf{z} | \mathbf{x}_c) || p(\mathbf{z} | \mathbf{x}_1, \dots, \mathbf{x}_C))$$

Evidence Lower bound (ELBO)

$$\frac{1}{C} \sum_{c=1}^C \mathbb{E}_{q(\mathbf{z} | \mathbf{x}_c)} \left[\sum_{i=1}^C \ln p(\mathbf{x}_i | \mathbf{z}) \right] - \mathcal{D}_{\text{KL}} (q(\mathbf{z} | \mathbf{x}_c) || p(\mathbf{z}))$$

Encoding for given channel
Reconstruction of all channels

Regularization: sparsity inducing prior

[Kingma et al, NIPS, 2015; Molchanov et al, ICML 2017]

Classic Implementation (non sparse)

Evidence Lower bound (ELBO)

$$\frac{1}{C} \sum_{c=1}^C \mathbb{E}_{q(\mathbf{z}|\mathbf{x}_c)} \left[\sum_{i=1}^C \ln p(\mathbf{x}_i|\mathbf{z}) \right] - \mathcal{D}_{\text{KL}}(q(\mathbf{z}|\mathbf{x}_c) || p(\mathbf{z}))$$

Reconstructions: $p(\mathbf{x}_i|\mathbf{z}) = \mathcal{N}(\mathbf{x}_i|\boldsymbol{\mu}_{\mathbf{z}}; \boldsymbol{\Sigma}_{\mathbf{z}})$

Encodings: $q(\mathbf{z}|\mathbf{x}_c) = \mathcal{N}(\mathbf{z}|\boldsymbol{\mu}_{\mathbf{x}_c}; \boldsymbol{\Sigma}_{\mathbf{x}_c})$

Prior: $p(\mathbf{z}) = \mathcal{N}(\mathbf{0}; \mathbf{I})$

Sparse Implementation

Evidence Lower bound (ELBO)

$$\frac{1}{C} \sum_{c=1}^C \mathbb{E}_{q(\mathbf{z}|\mathbf{x}_c)} \left[\sum_{i=1}^C \ln p(\mathbf{x}_i|\mathbf{z}) \right] - \mathcal{D}_{\text{KL}}(q(\mathbf{z}|\mathbf{x}_c) || p(\mathbf{z}))$$

Reconstructions: $p(\mathbf{x}_i|\mathbf{z}) = \text{same}$

Encodings: $q(\mathbf{z}|\mathbf{x}_c) = \mathcal{N}(\mathbf{z}|\boldsymbol{\mu}_{\mathbf{x}_c}; \boldsymbol{\alpha} \odot \boldsymbol{\mu}_{\mathbf{x}_c}^2)$

Prior: $p(\mathbf{z}) \propto 1/|\mathbf{z}|$

Sparse Implementation: why it works?

Evidence Lower bound (ELBO)

$$\frac{1}{C} \sum_{c=1}^C \mathbb{E}_{q(\mathbf{z}|\mathbf{x}_c)} \left[\sum_{i=1}^C \ln p(\mathbf{x}_i|\mathbf{z}) \right] - \mathcal{D}_{\text{KL}}(q(\mathbf{z}|\mathbf{x}_c) || p(\mathbf{z}))$$

Encodings: $q(\mathbf{z}|\mathbf{x}_c) = \mathcal{N}(\mathbf{z}|\boldsymbol{\mu}_{\mathbf{x}_c}; \boldsymbol{\alpha} \odot \boldsymbol{\mu}_{\mathbf{x}_c}^2)$

Prior: $p(\mathbf{z}) \propto 1/|\mathbf{z}|$

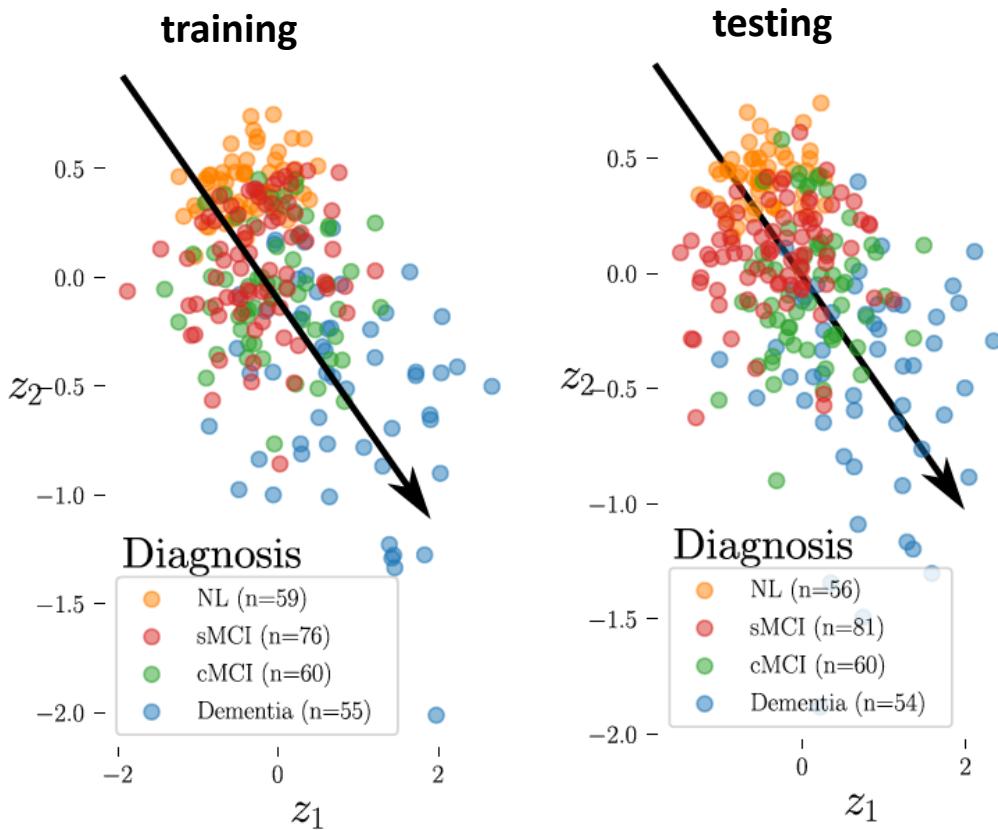
Prior and encodings act together such that (element-wise):

$$\lim_{\mu_i \rightarrow 0} \mathcal{N}(z_i | \mu_i; \alpha_i \cdot \mu_i^2) = \delta(0)$$

Relationship between α and the probability of pruning the i -th dimension:

$$\alpha_i = \frac{p_i}{1 - p_i}$$

Prediction from latent space

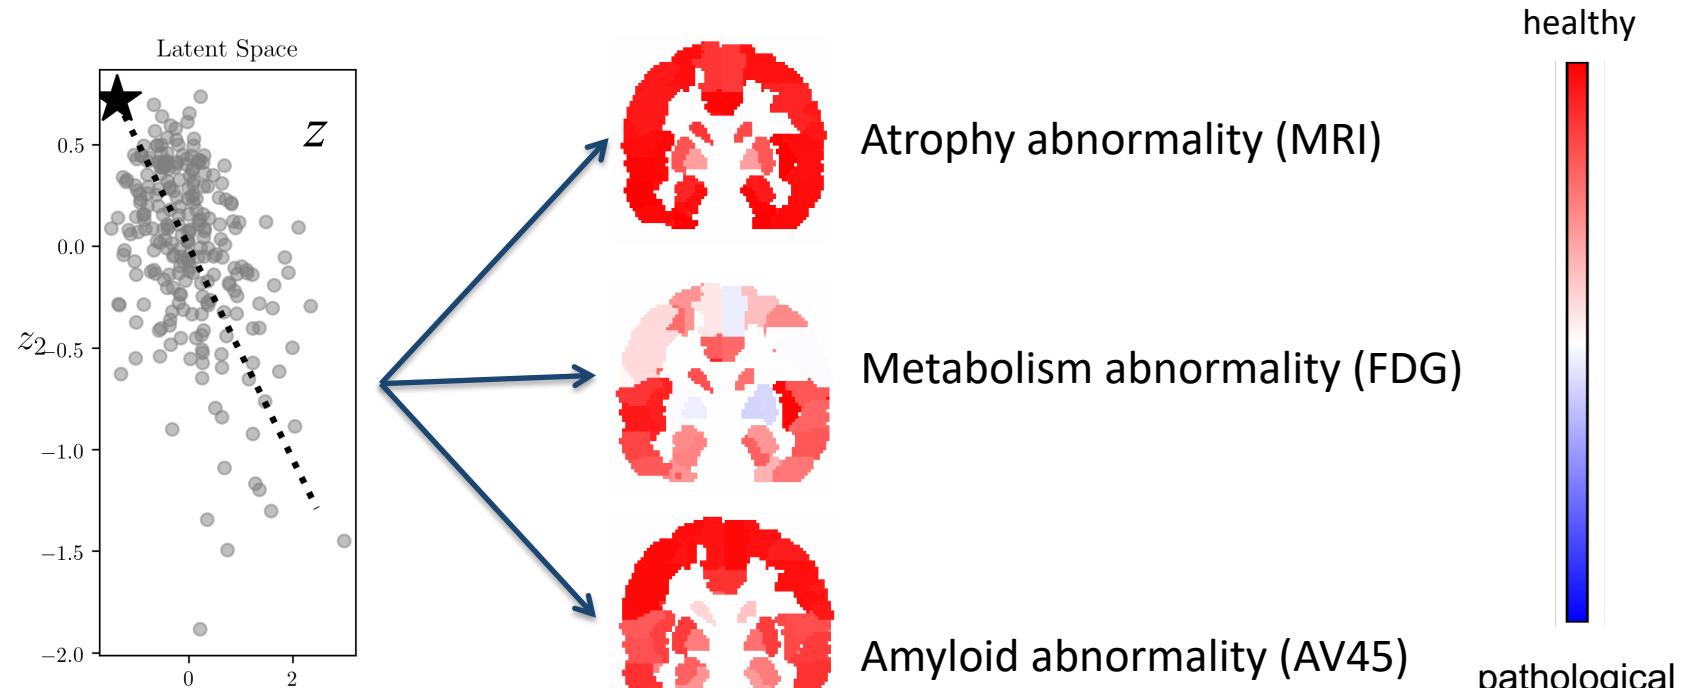


Joint modeling of

- Brain imaging:
 - Structural (T1 MRI)
 - Molecular (FDG-PET + Amy-PET)
- Socio-demographic factors
- Clinical scores

	Accuracy (SD)
Cognitively Healthy	0.89 (0.03)
Stable Mild Cognitive Impairment (sMCI)	0.75 (0.02)
MCI converting to Dementia (cMCI)	0.70 (0.05)
Dementia	0.94 (0.05)

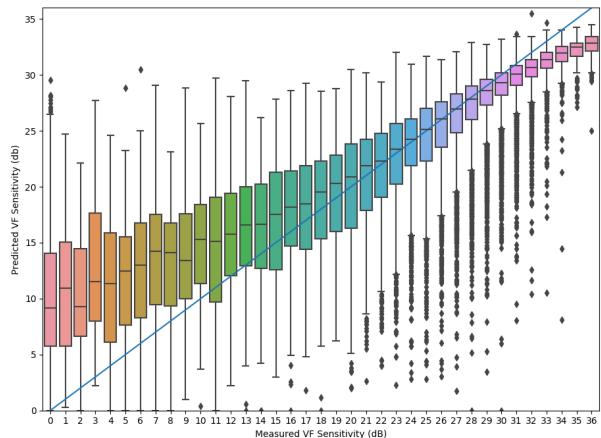
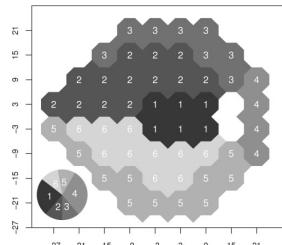
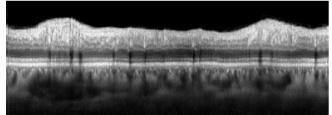
Generation from latent space



- Improved interpretability
- Multi-channel: working with missing data/data imputation
- Simulations for clinical trials

Large-scale applications

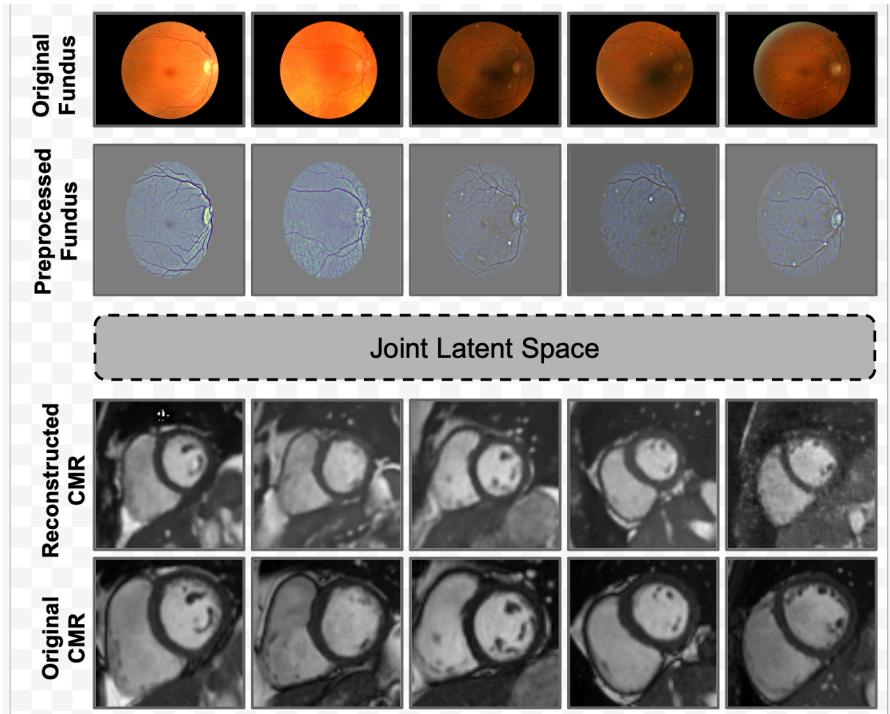
From RNFL thickness to visual fields



Courtesy of Lazaridis et al.

work in progress, UCL-Inria collaboration

From fundus to cardiac images



Courtesy of Diaz-Pinto et al.

work in progress at CISTIB, University of Leeds, UK

Thank you